

Power Management from Linux Kernel to Android

Matt Hsu & Jim Huang (jserv)
from Oxlab

June 19, 2009





Rights to copy

© Copyright 2009 **0xlab.org**
contact@0xlab.org

Corrections, suggestions, contributions and translations are welcome!



Attribution – ShareAlike 3.0

You are free

- to copy, distribute, display, and perform the work
- to make derivative works
- to make commercial use of the work

Under the following conditions

- **BY: Attribution.** You must give the original author credit.
- **Share Alike.** If you alter, transform, or build upon this work, you may distribute the resulting work only under a license identical to this one.
 - For any reuse or distribution, you must make clear to others the license terms of this work.
 - Any of these conditions can be waived if you get permission from the copyright holder.

Your fair use and other rights are in no way affected by the above.

License text: <http://creativecommons.org/licenses/by-sa/3.0/legalcode>



Agenda

- ▶ Introduction to Linux Power Management
- ▶ Concepts behind Android Power Management
- ▶ Design and Implementation
- ▶ Room for Improvements



- ▶ **Introduction to Linux Power Management**
 - ▶ **Goal of Power Management**
 - ▶ **APM vs. ACPI**
 - ▶ **APM emulation**
 - ▶ **No silver bullet**
 - ▶ **Manage power in different power state**
 - ▶ **Lighthouse in the sea**
- ▶ Concepts behind Android Power Management
- ▶ Design and Implementation
- ▶ Room for Improvements

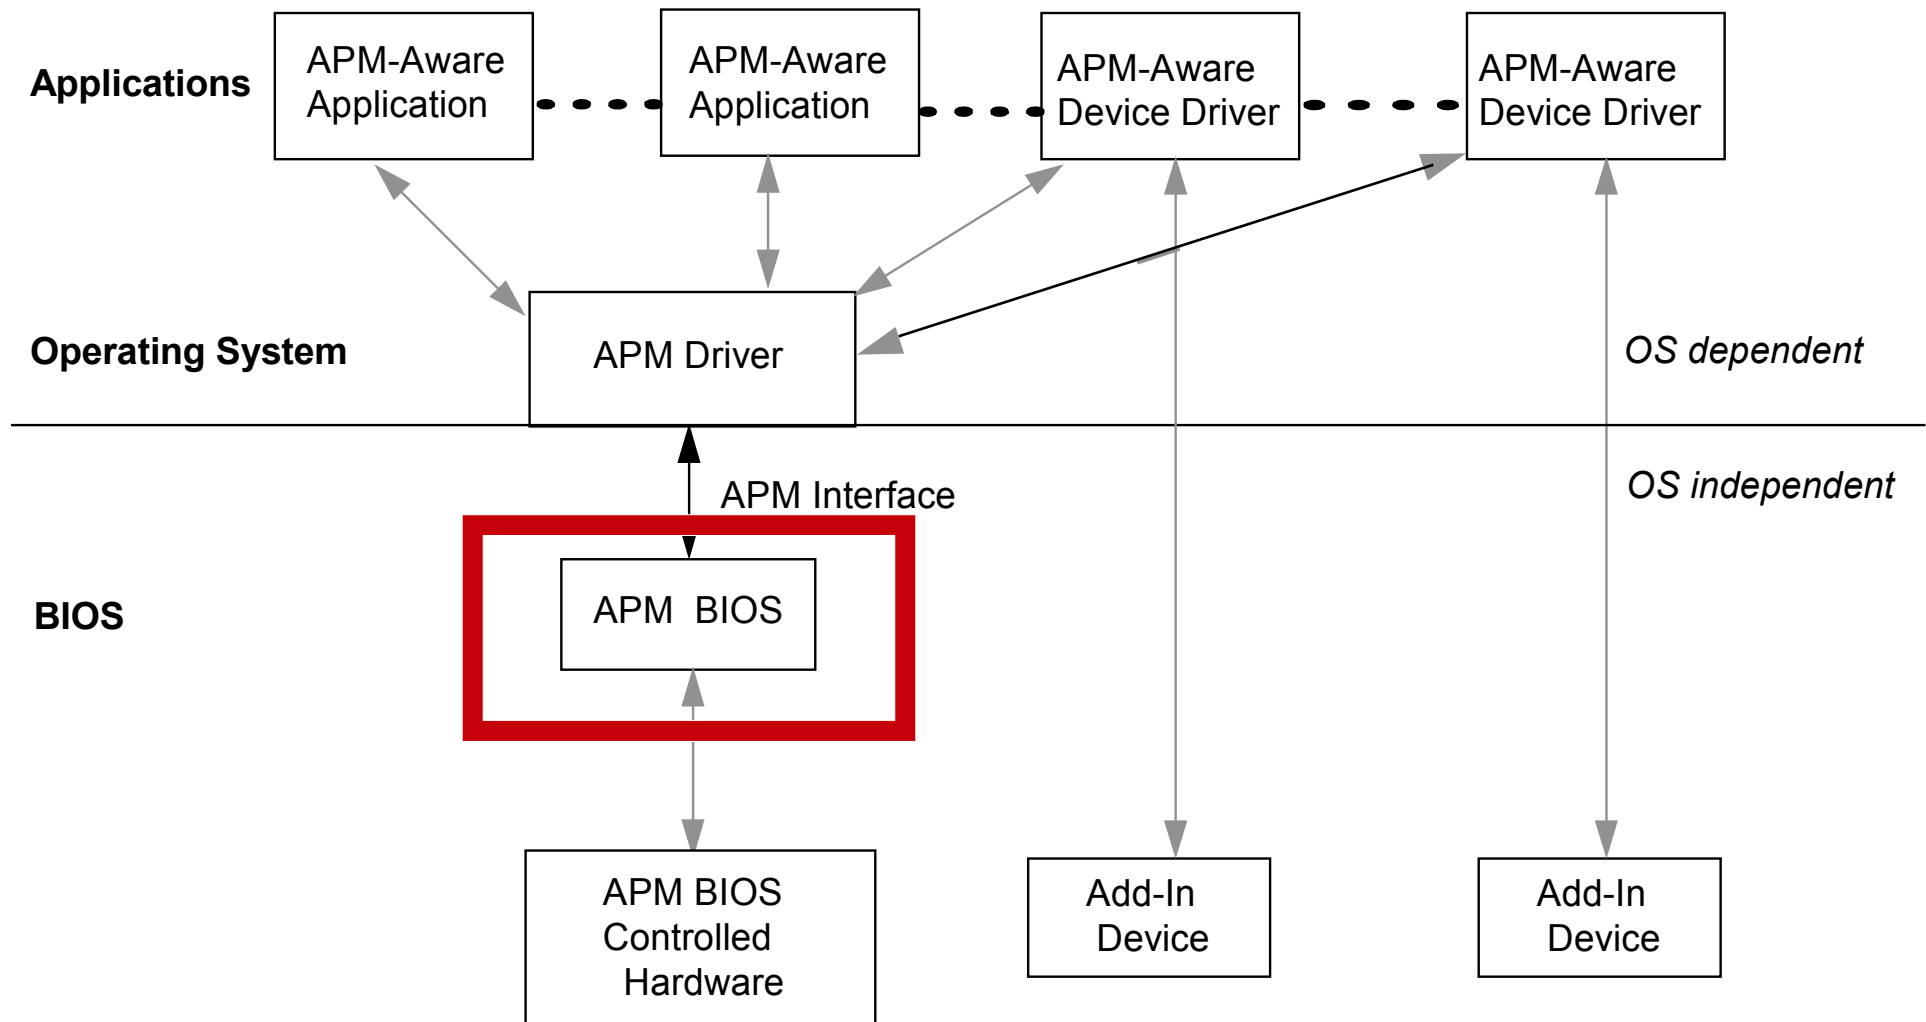
Prolong



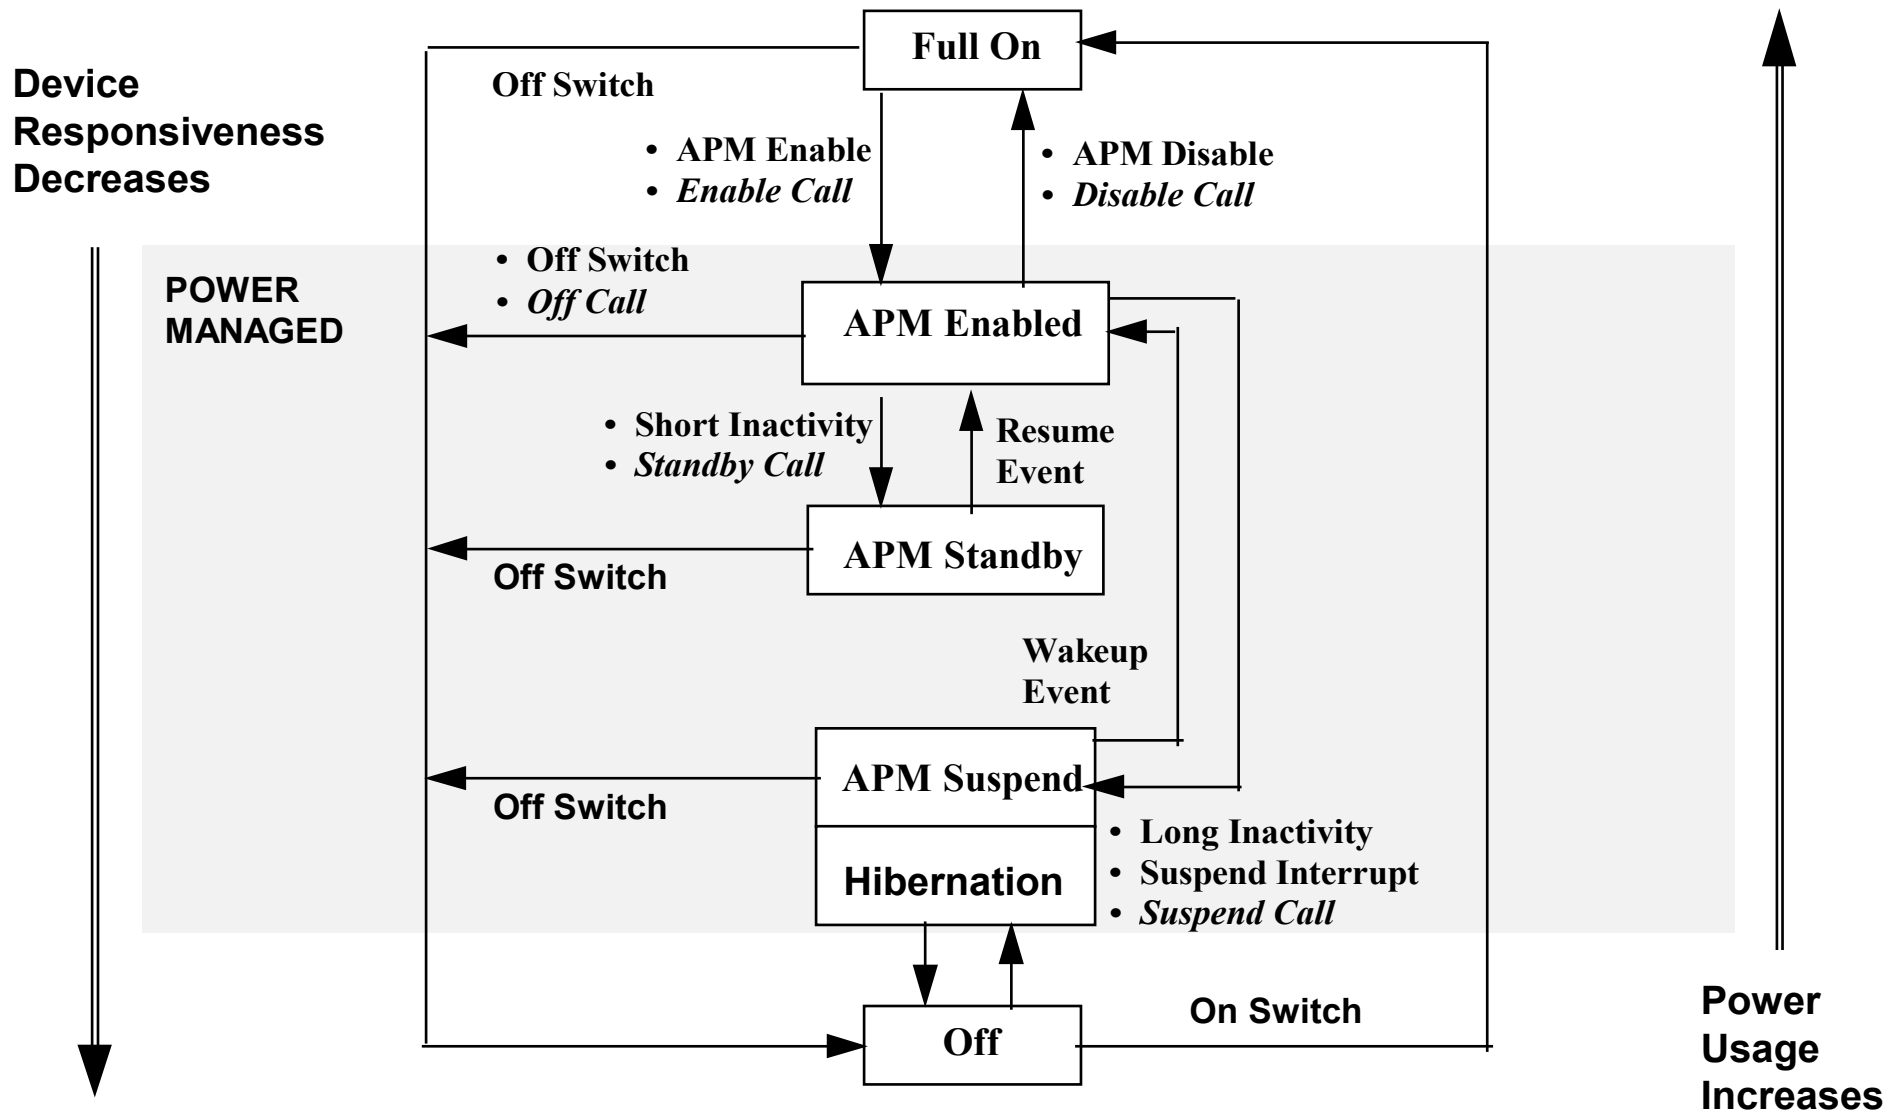
Life

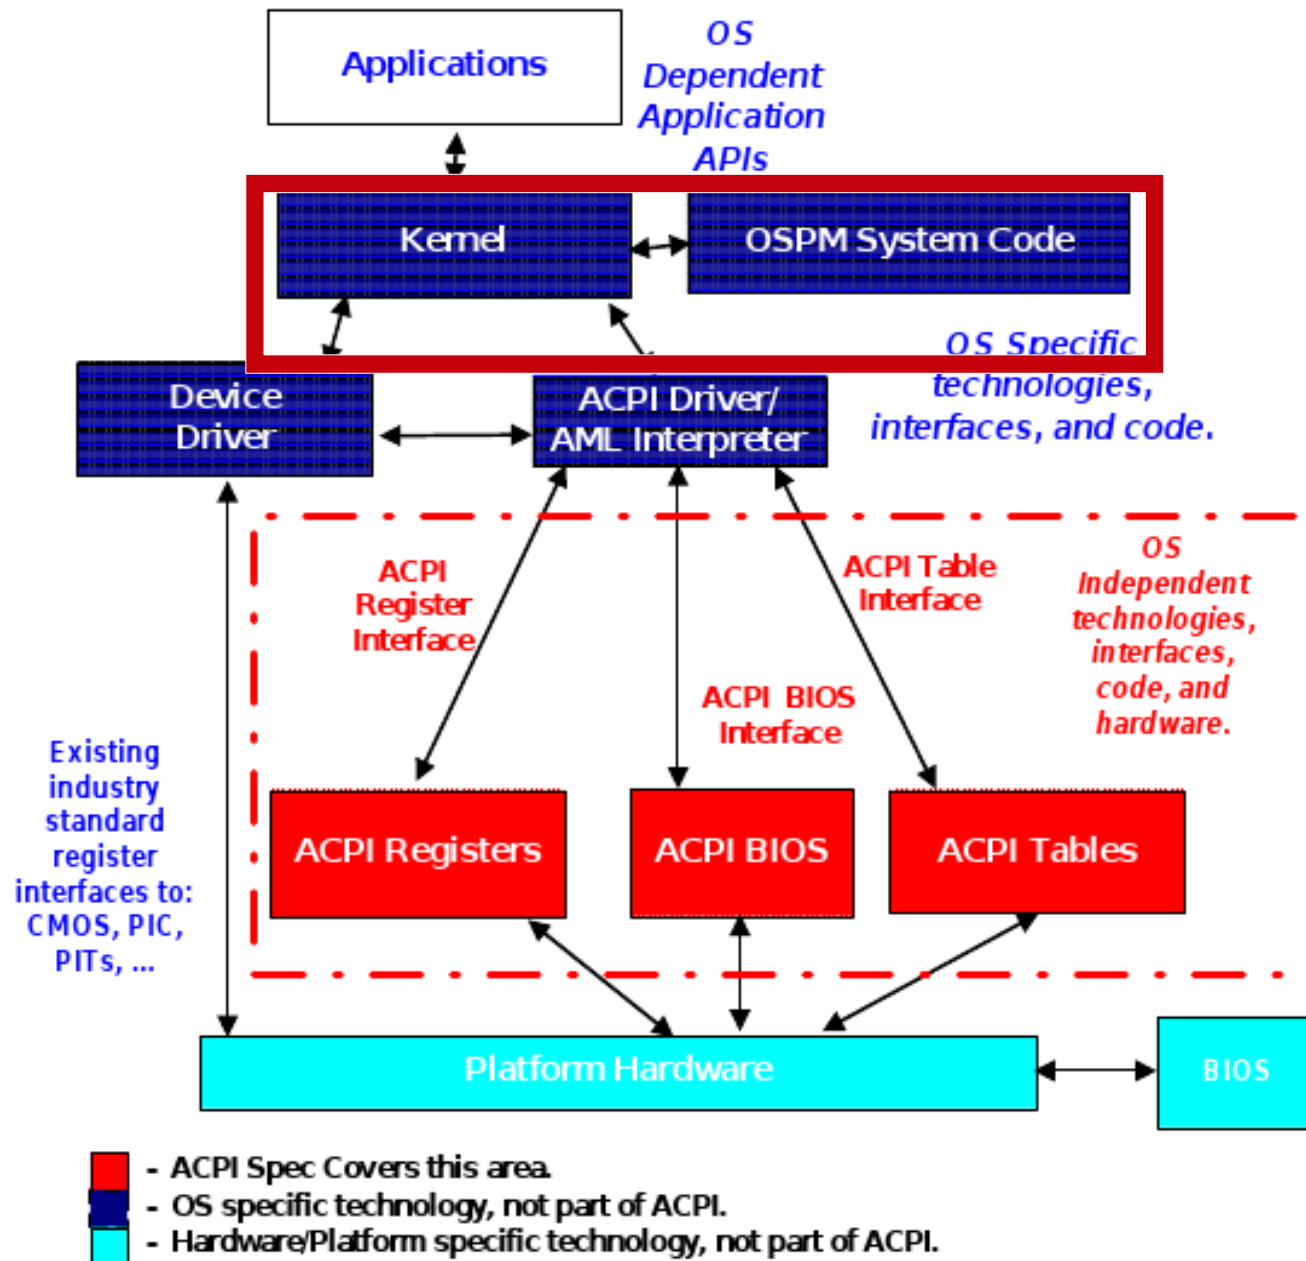


Advanced Power Management

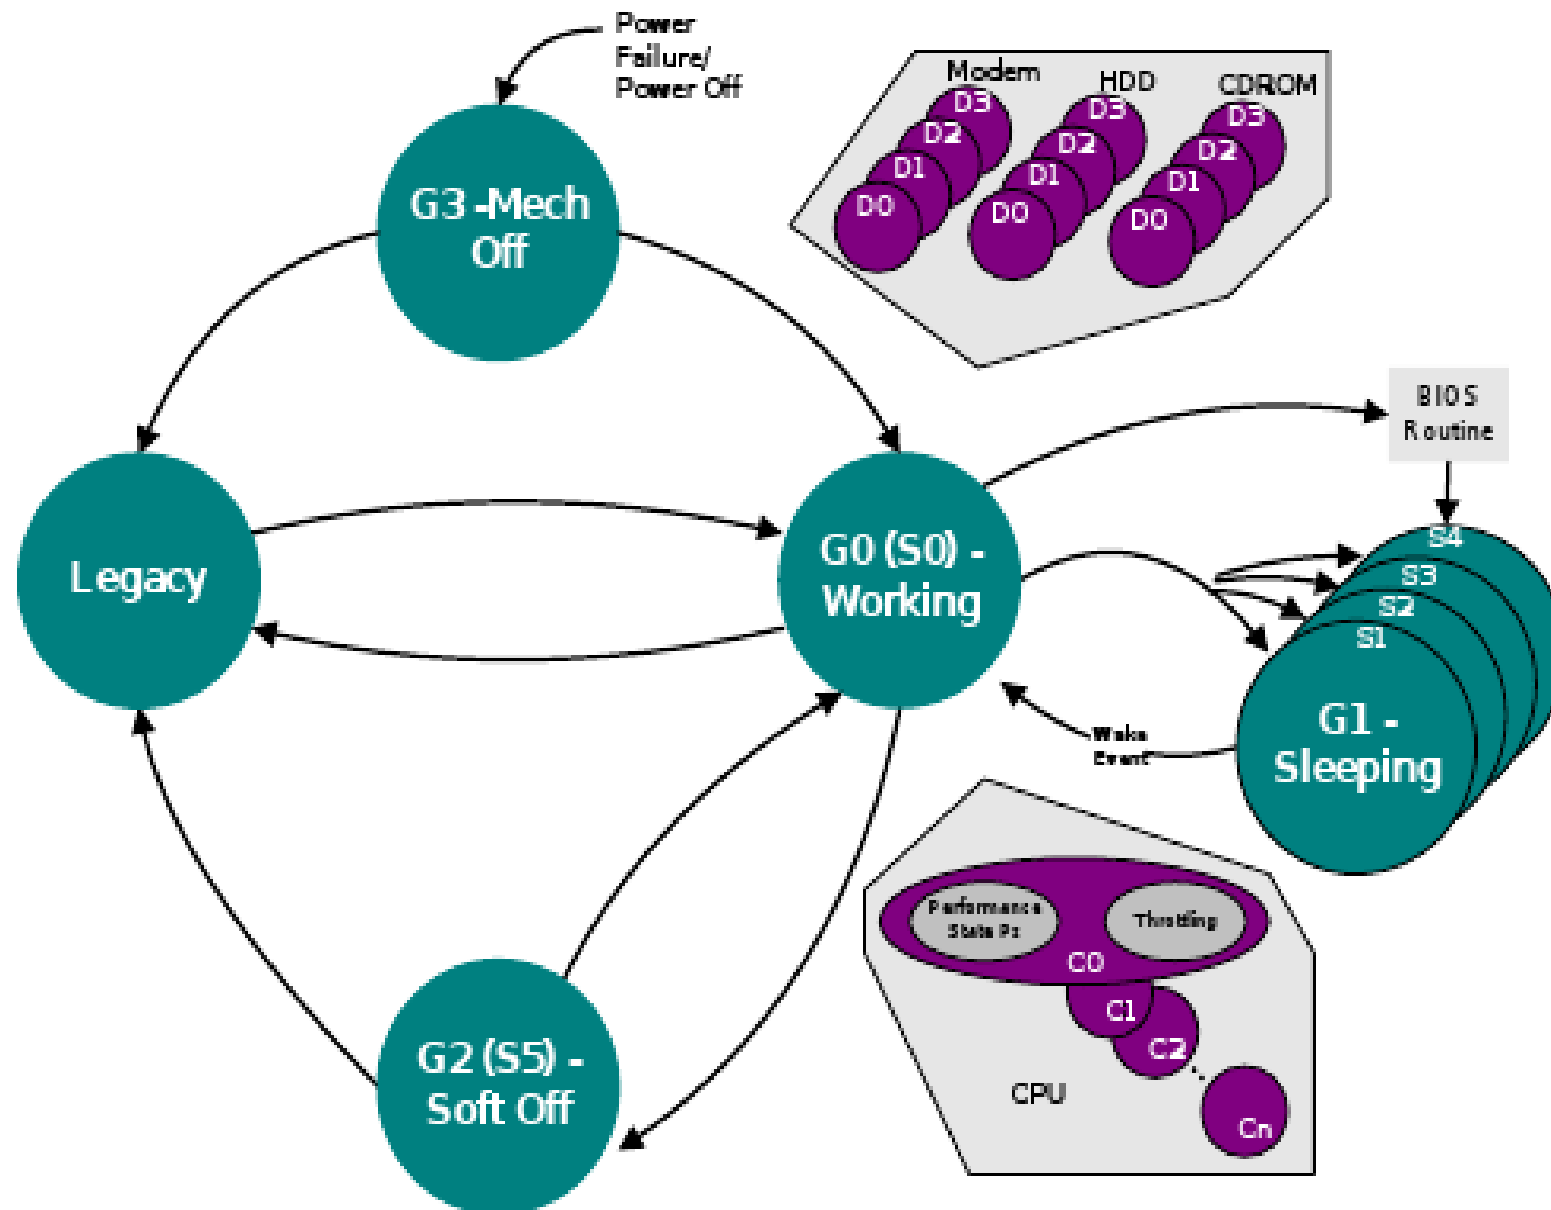


APM power state





ACPI Power State Transition





APM vs. ACPI

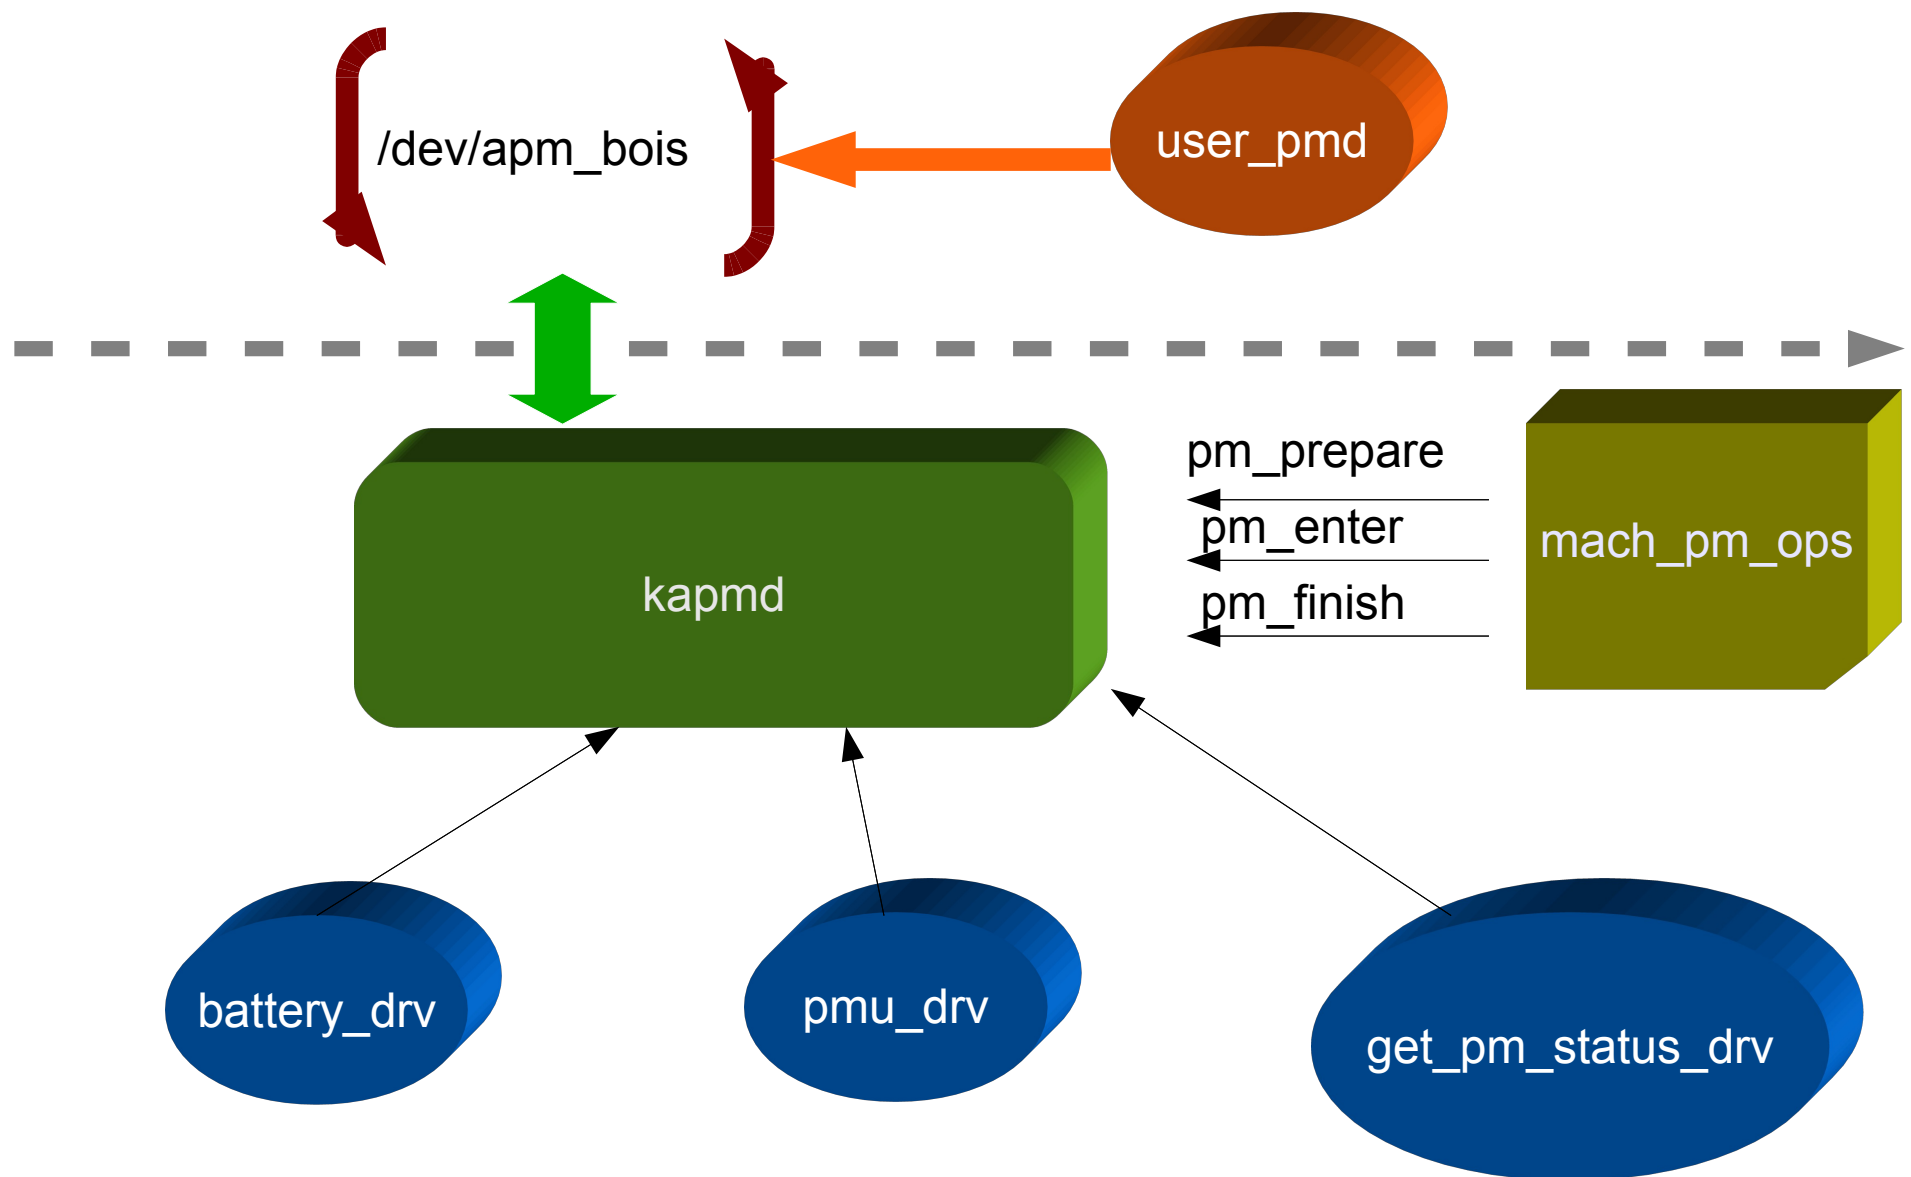
APM

- ▶ Control resides in BIOS
- ▶ Uses activity timeouts to determine when to power down a device
- ▶ BIOS rarely used in embedded systems
- ▶ Makes power-management decisions without informing OS or individual applications
- ▶ No knowledge of add-in cards or new devices (e.g. USB, IEEE 1394)

ACPI

- ▶ Control divided between BIOS and OS
- ▶ Decisions managed through the OS
- ▶ Enables sophisticated power policies for general-purpose computers with standard usage patterns and hardware
- ▶ No knowledge of device-specific scenarios (e.g. need to provide predictable response times or to respond to critical events over extended period)

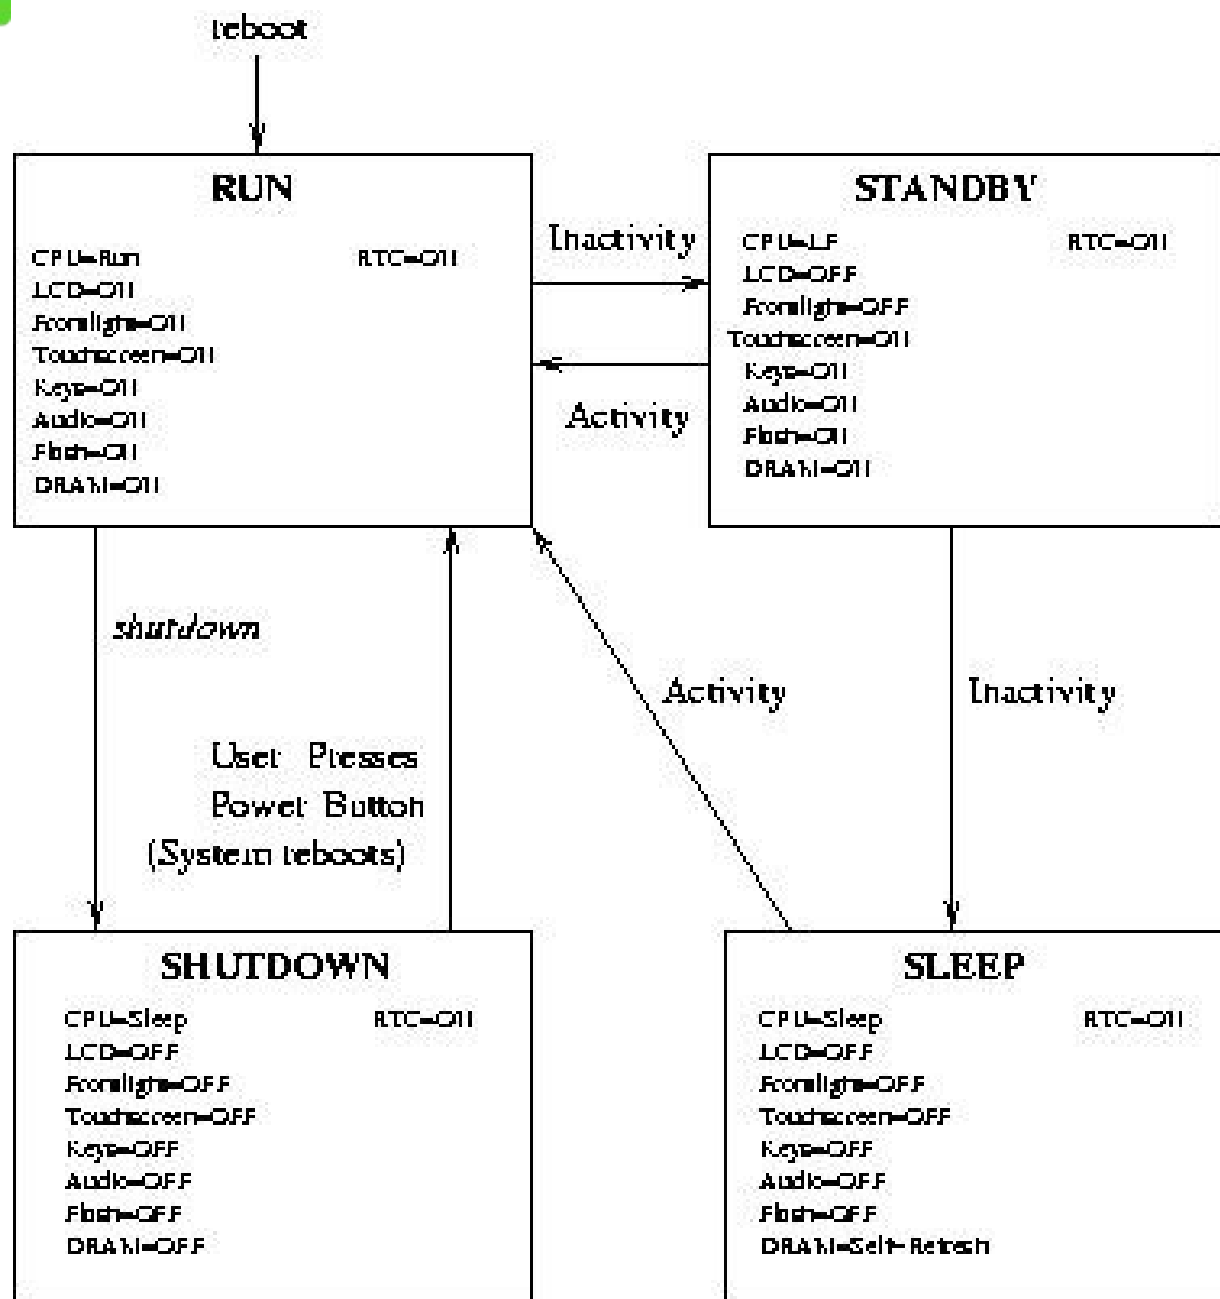
APM emulation





No silver bullet
for building most power-saving
mobile system.

Power State Transition





Runtime / Standby

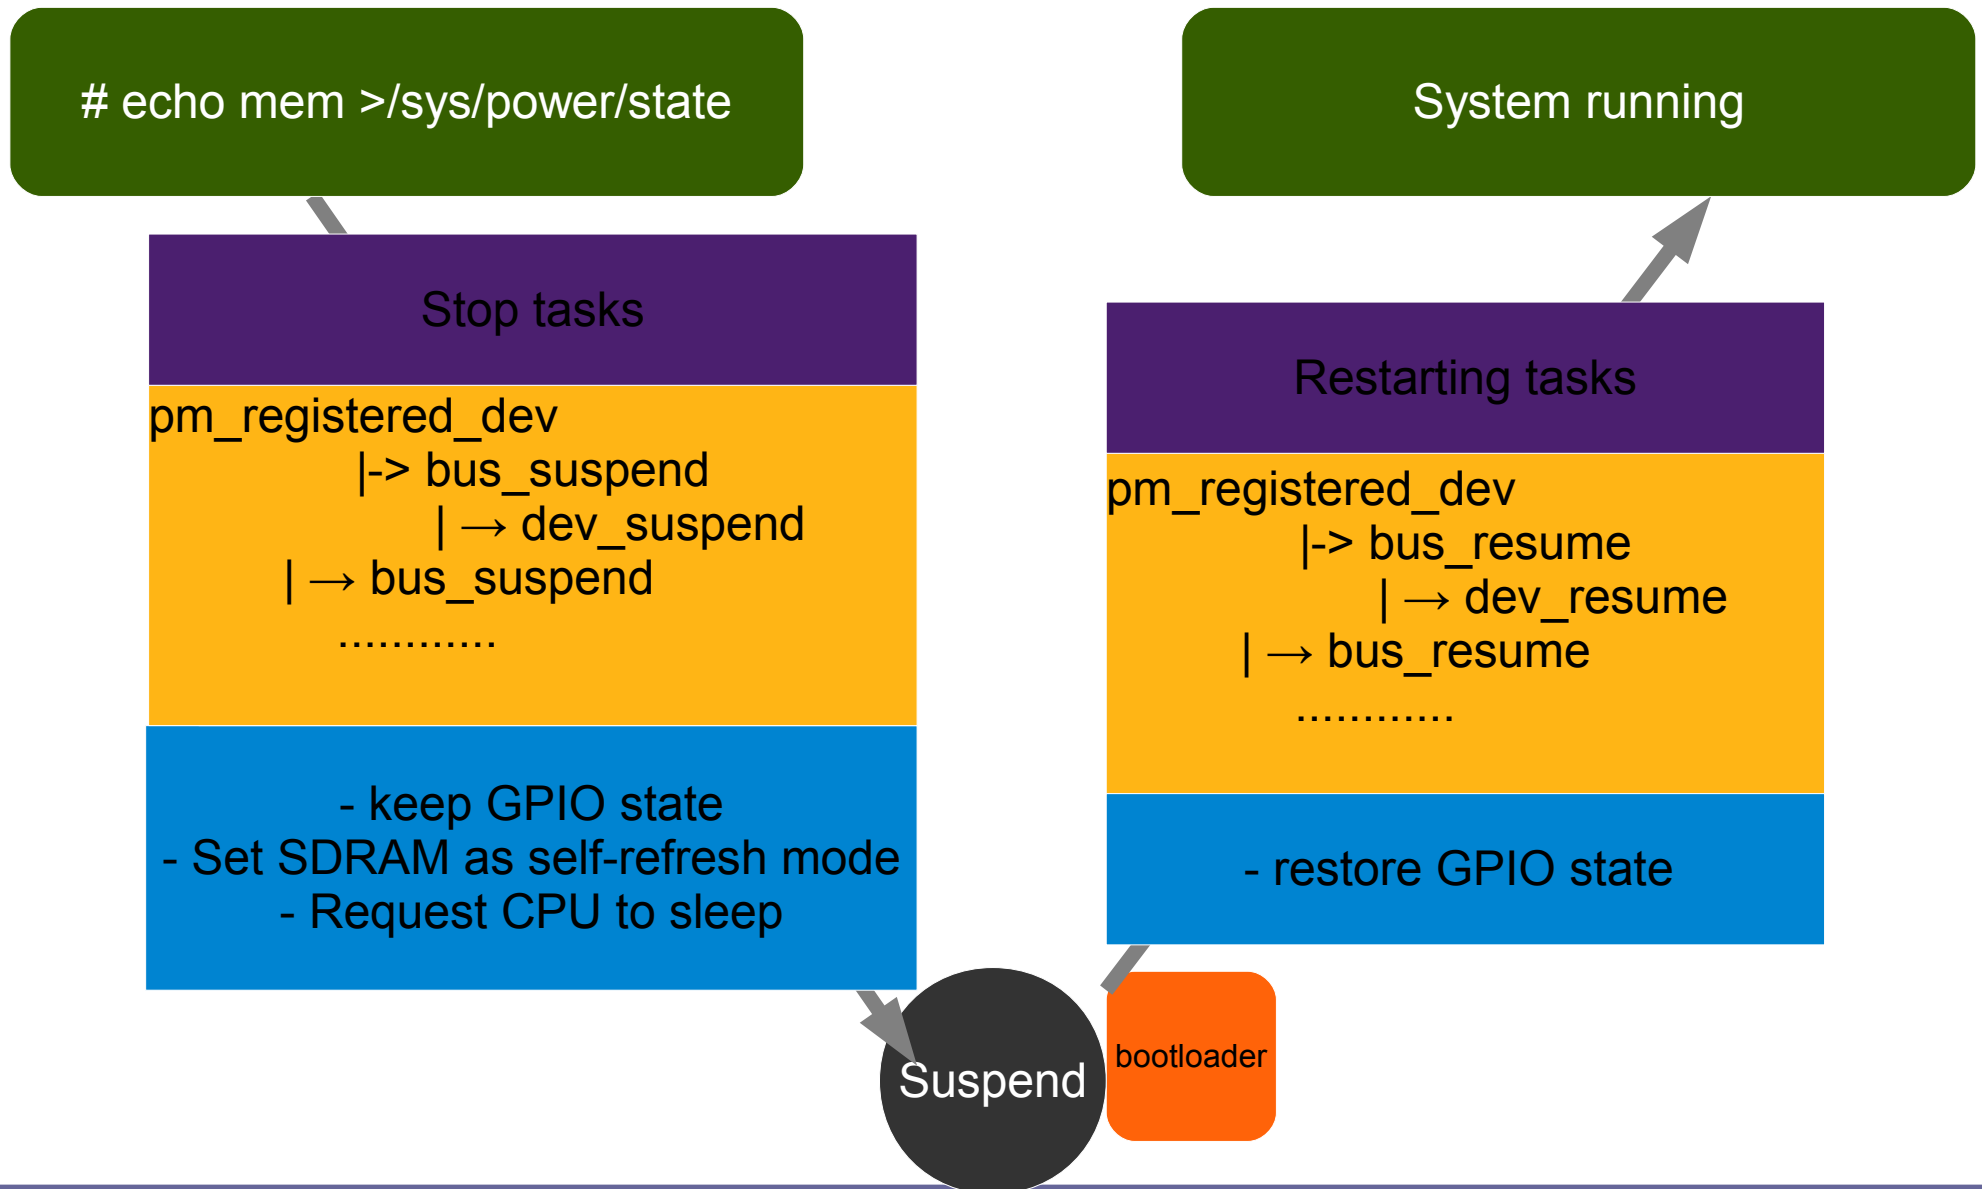
- ▶ Application-driven power management
- ▶ Micro manage your device
 - ▶ WiFi -enable PS poll mode
 - ▶ Switch on/off device by demand (Android's concept)
 - ▶ Gating off unused device clock
- ▶ Keep the flexibility of CPU
 - ▶ Gating CPU freq dynamically
 - ▶ Tickless idle
 - ▶ Using DMA instead PIO

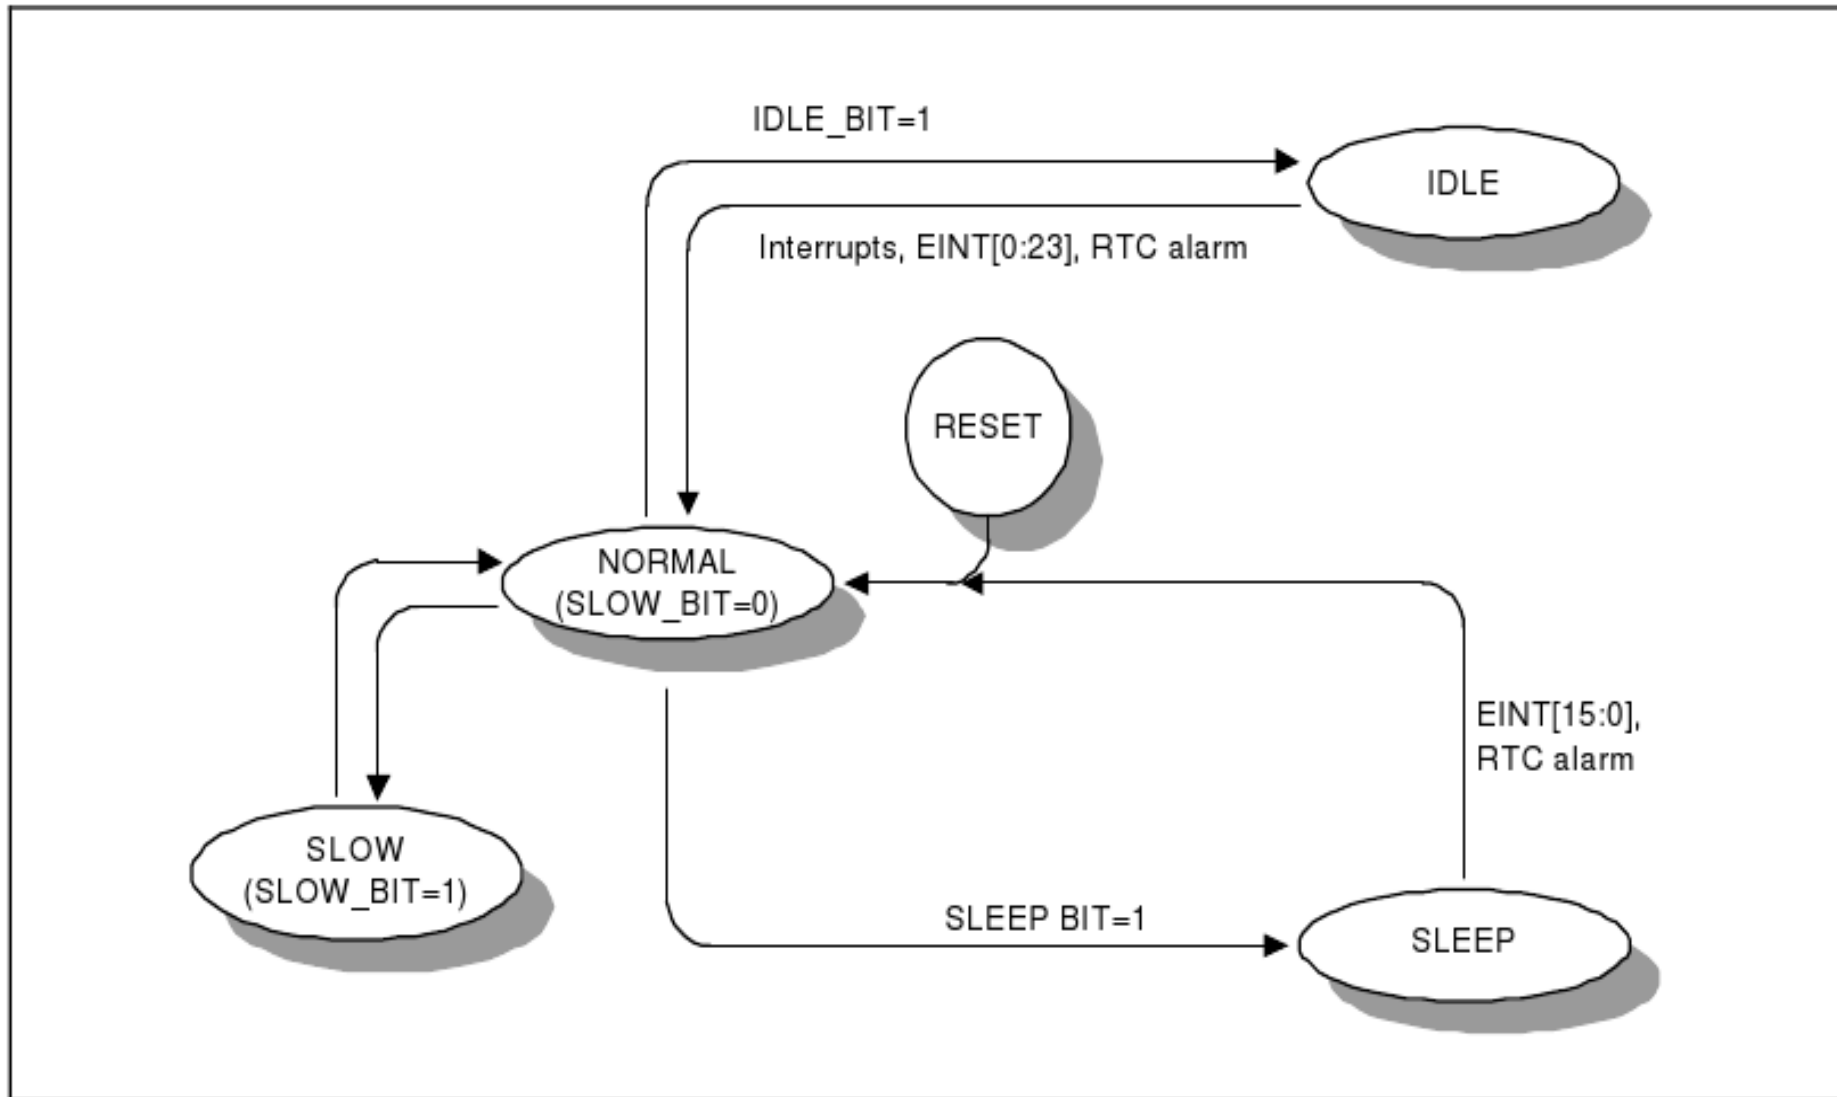


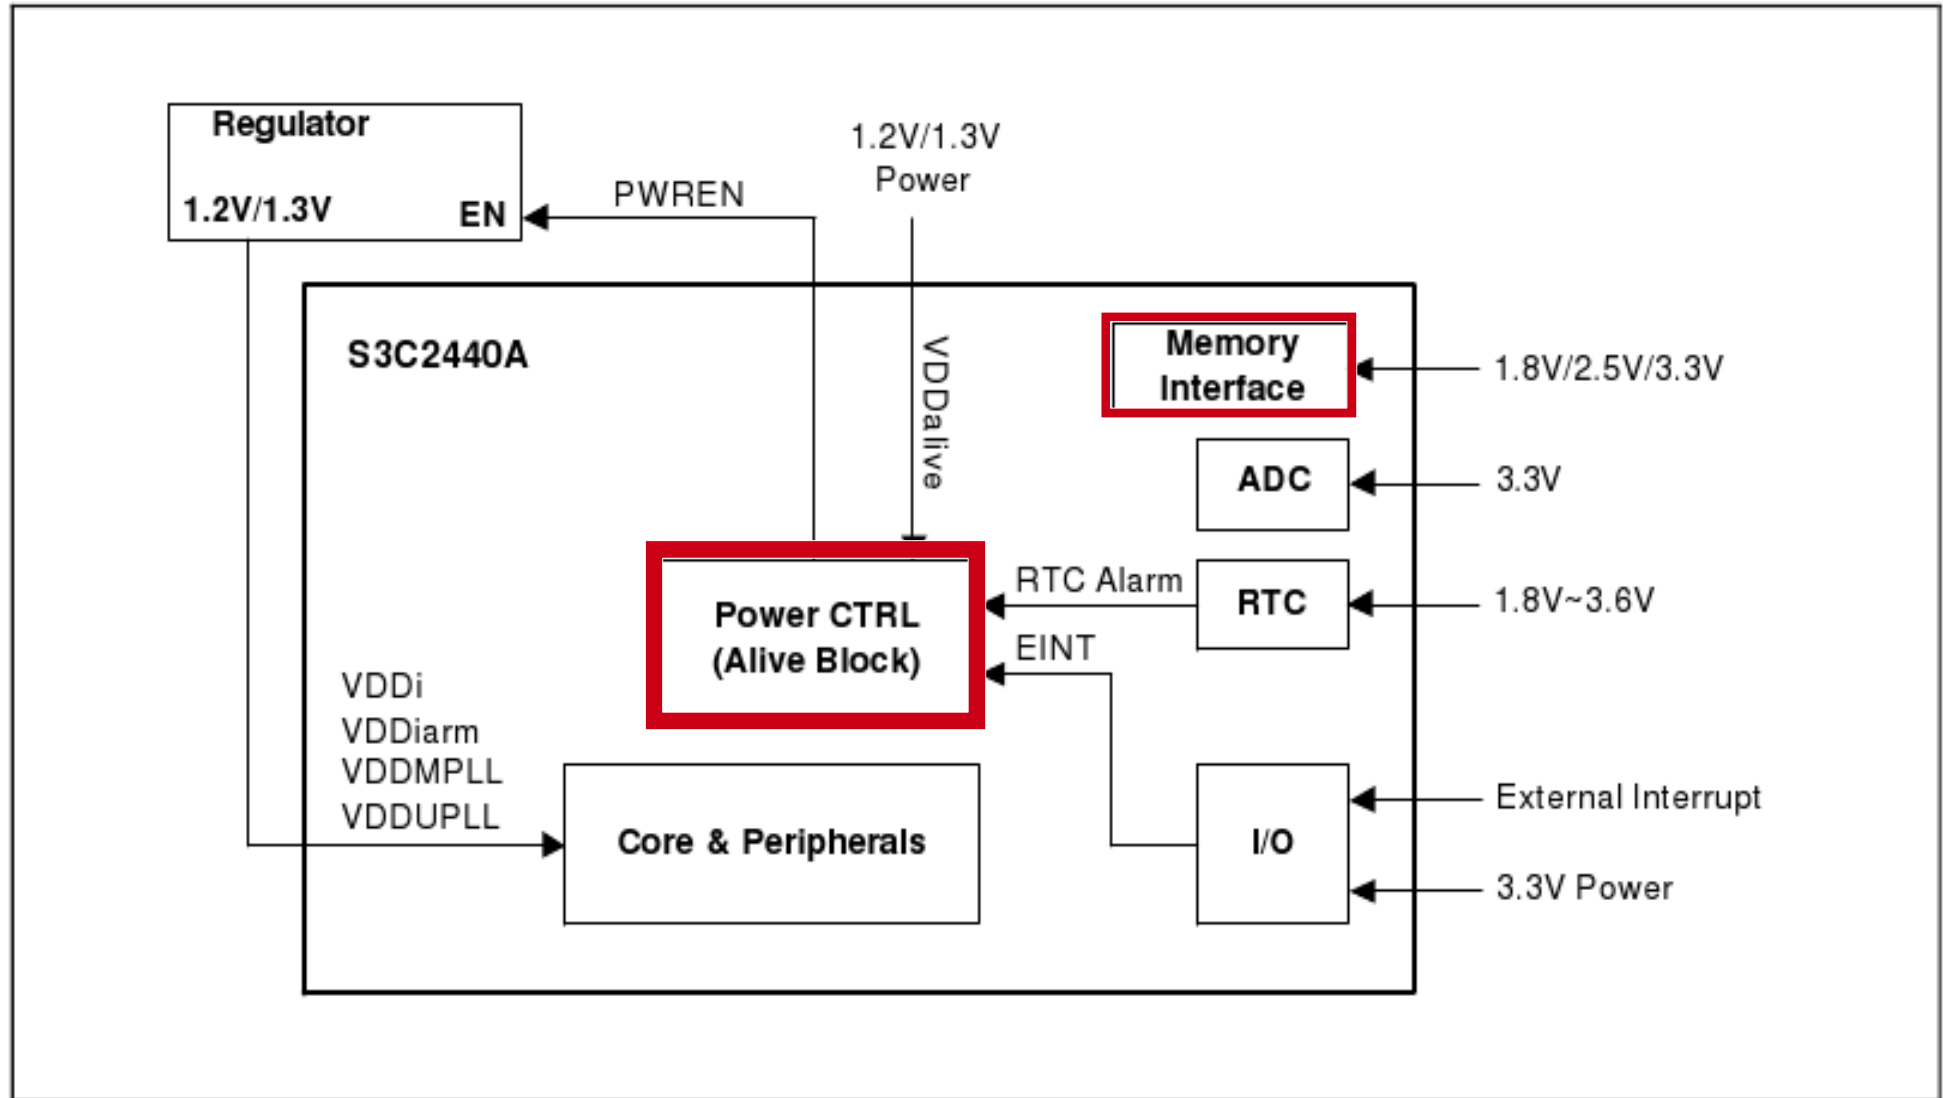
Suspend/Sleep Mode

- ▶ STR vs STD
 - ▶ Suspend to RAM
 - ▶ Suspend to Disk/Hibernate
- ▶ STR: A technique by which systems preserve state in RAM during suspend and restore system state from RAM upon resume
- ▶ STR is suitable for mobile device if hardware supports this function

Suspend/Resume Control Path









A ship on the beach is a lighthouse to the sea

- ▶ Provide a kernel module to dump and check GPIOs
- ▶ Provide power source as independent as possible
- ▶ Don't ignore the poweroff state
- ▶ To most difficult part of debugging STR is resuming

- ▶ Introduction to Linux Power Management
- ▶ Concepts behind Android Power Management
 - ▶ Basically, it is the slim wrapper for Linux Power Management
- ▶ Design and Implementation
- ▶ Room for Improvements



Where is Power Manager?




PM?



In fact, power ghost exists everywhere

- ▶ Let “**grep -ri power**” tell you about the details!
- ▶ It is layered into several components, but implementation involves in some hacks.
- ▶ Check: CONFIG_PM, sysfs, device drivers, libhardware_legacy, libril, init.rc, powerd, alarm, vold, JNI, PowerManager, BatteryManager, ...





Base: Linux Kernel

Directory: arch/arm/mach-omap2/
board-omap3beagle.c
gpmc.c
pm.[ch]



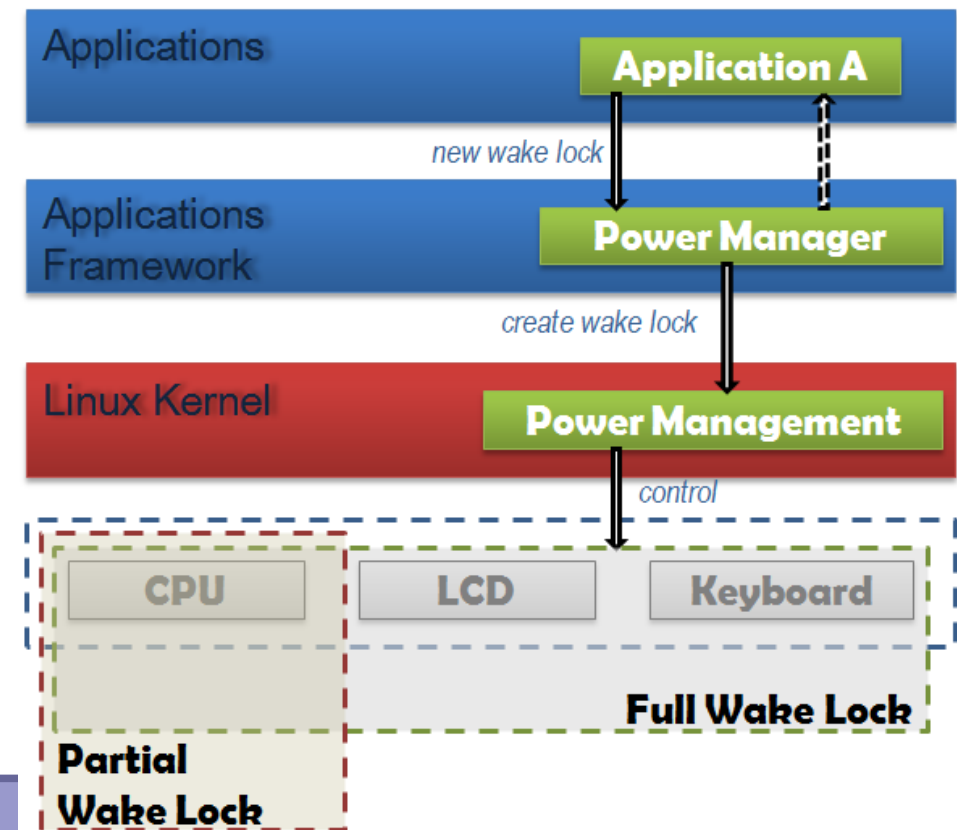
- ▶ Android does rely on Linux Kernel 2.6 for core system services
 - ▶ Memory/Process Management
 - ▶ Device Driver Model
 - ▶ sysfs, kobject/uevent, netlink
- ▶ Android Kernel extensions
 - ▶ Binder
 - ▶ android_power
 - ▶ /sys/android_power/, /sys/power/

Key Idea: Android attempts to provide an abstraction layer between hardware and the related software stack.



Android's PM Concepts

- ▶ Android PM is built on top of standard Linux Power Management.
- ▶ It can support more aggressive PM, but looks fairly simple now.
- ▶ Components make requests to keep the power on through **“Wake Locks”**.
 - ▶ PM does support several types of “Wake Locks”.



- ▶ If there are no active wake locks, CPU will be turned off.
- ▶ If there is are partial wake locks, screen and keyboard will be turned off.



Applications

Application A

Application B

Application C

```
Wl = newWakeLock(...);  
Wl.acquire();  
Wl.release();
```

Applications Framework

PowerManager

Android.os.PowerManager

Power

Android.os.Power

PowerManagerService

Android.server.PowerManagerService

JNI

Libraries (user space)

Core Libraries



Power

/lib/hardware/power.c

Linux Kernel

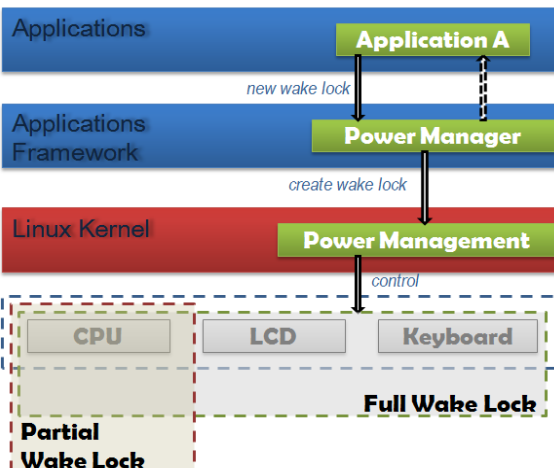
Linux Drivers

Android Power Management

/drivers/android/power.c

```
Android_register_early_suspend()  
Android_register_early_resume()
```

Linux Power Management





Types of Wake Locks

▶ ACQUIRE_CAUSES_WAKEUP

- ▶ Normally wake locks don't actually wake the device, they just cause it to remain on once it's already on. Think of the video player app as the normal behavior.

▶ FULL_WAKE_LOCK

- ▶ Wake lock that ensures that the screen and keyboard are on at full brightness.

▶ ON_AFTER_RELEASE

- ▶ When this wake lock is released, poke the user activity timer so the screen stays on for a little longer.

▶ PARTIAL_WAKE_LOCK

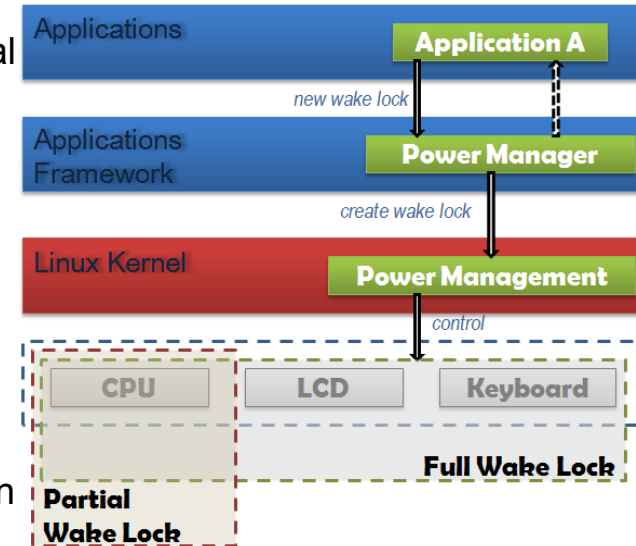
- ▶ Wake lock that ensures that the CPU is running. The screen might not be on.

▶ SCREEN_BRIGHT_WAKE_LOCK

- ▶ Wake lock that ensures that the screen is on at full brightness; the keyboard backlight will be allowed to go off.

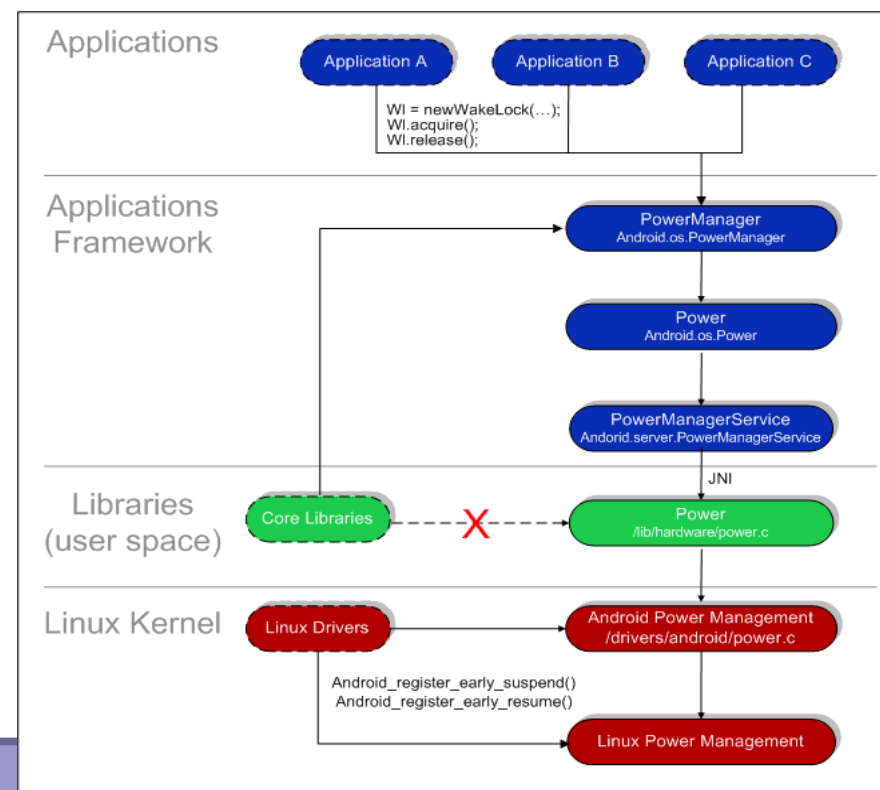
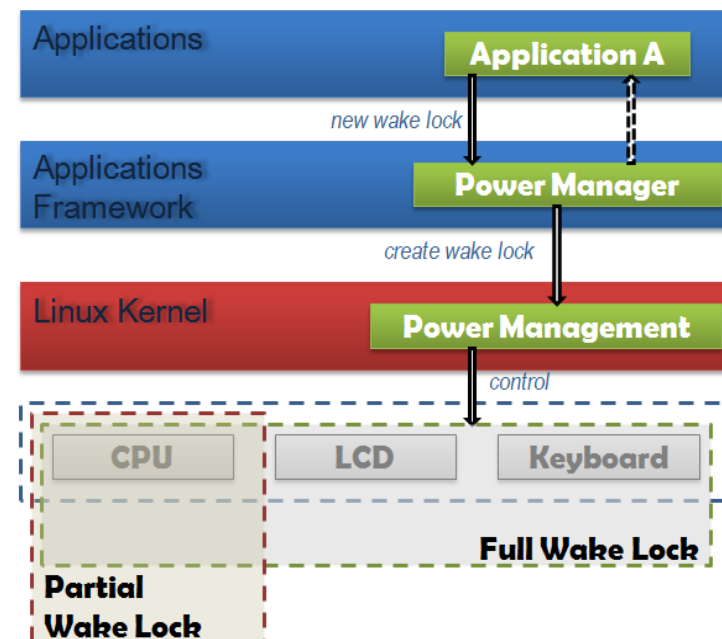
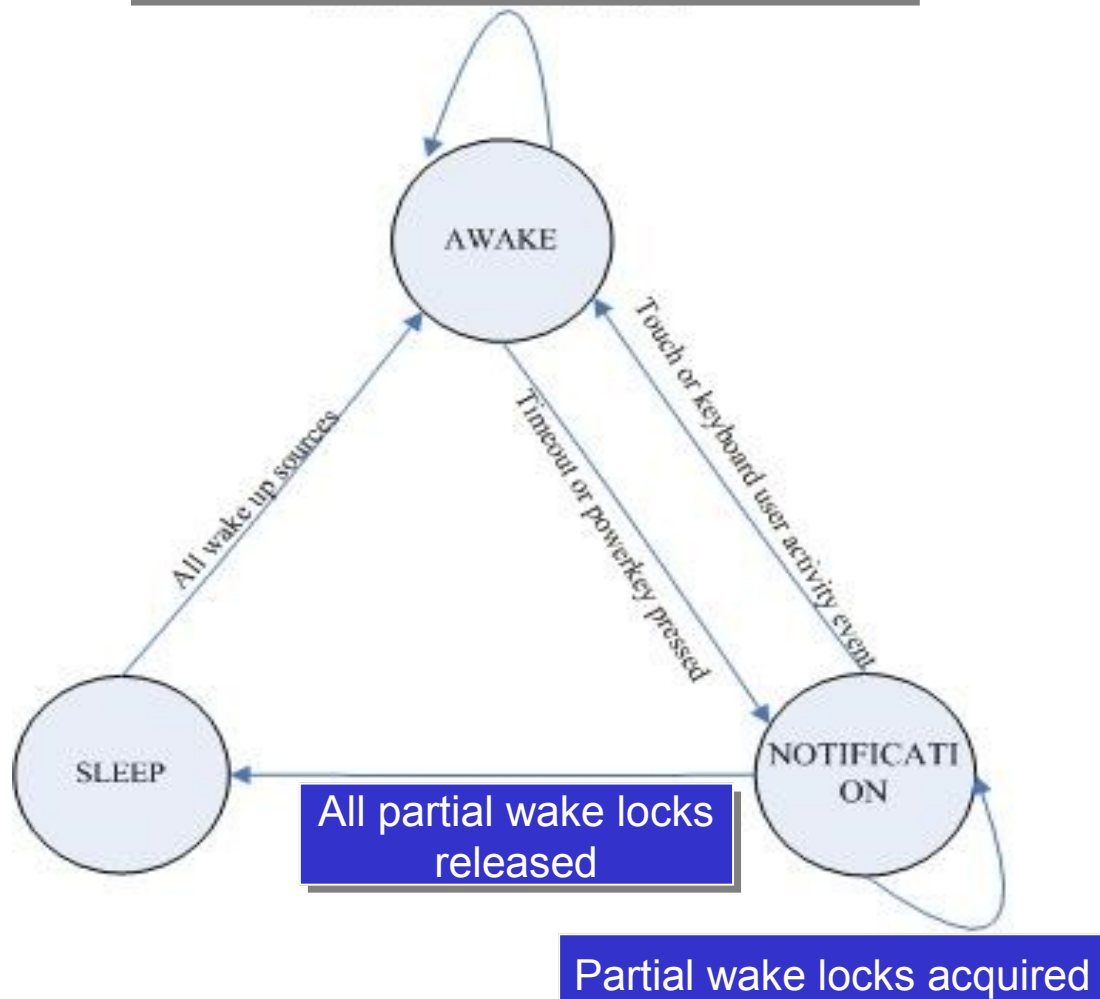
▶ SCREEN_DIM_WAKE_LOCK

- ▶ Wake lock that ensures that the screen is on, but the keyboard backlight will be allowed to go off, and the screen backlight will be allowed to go dim.



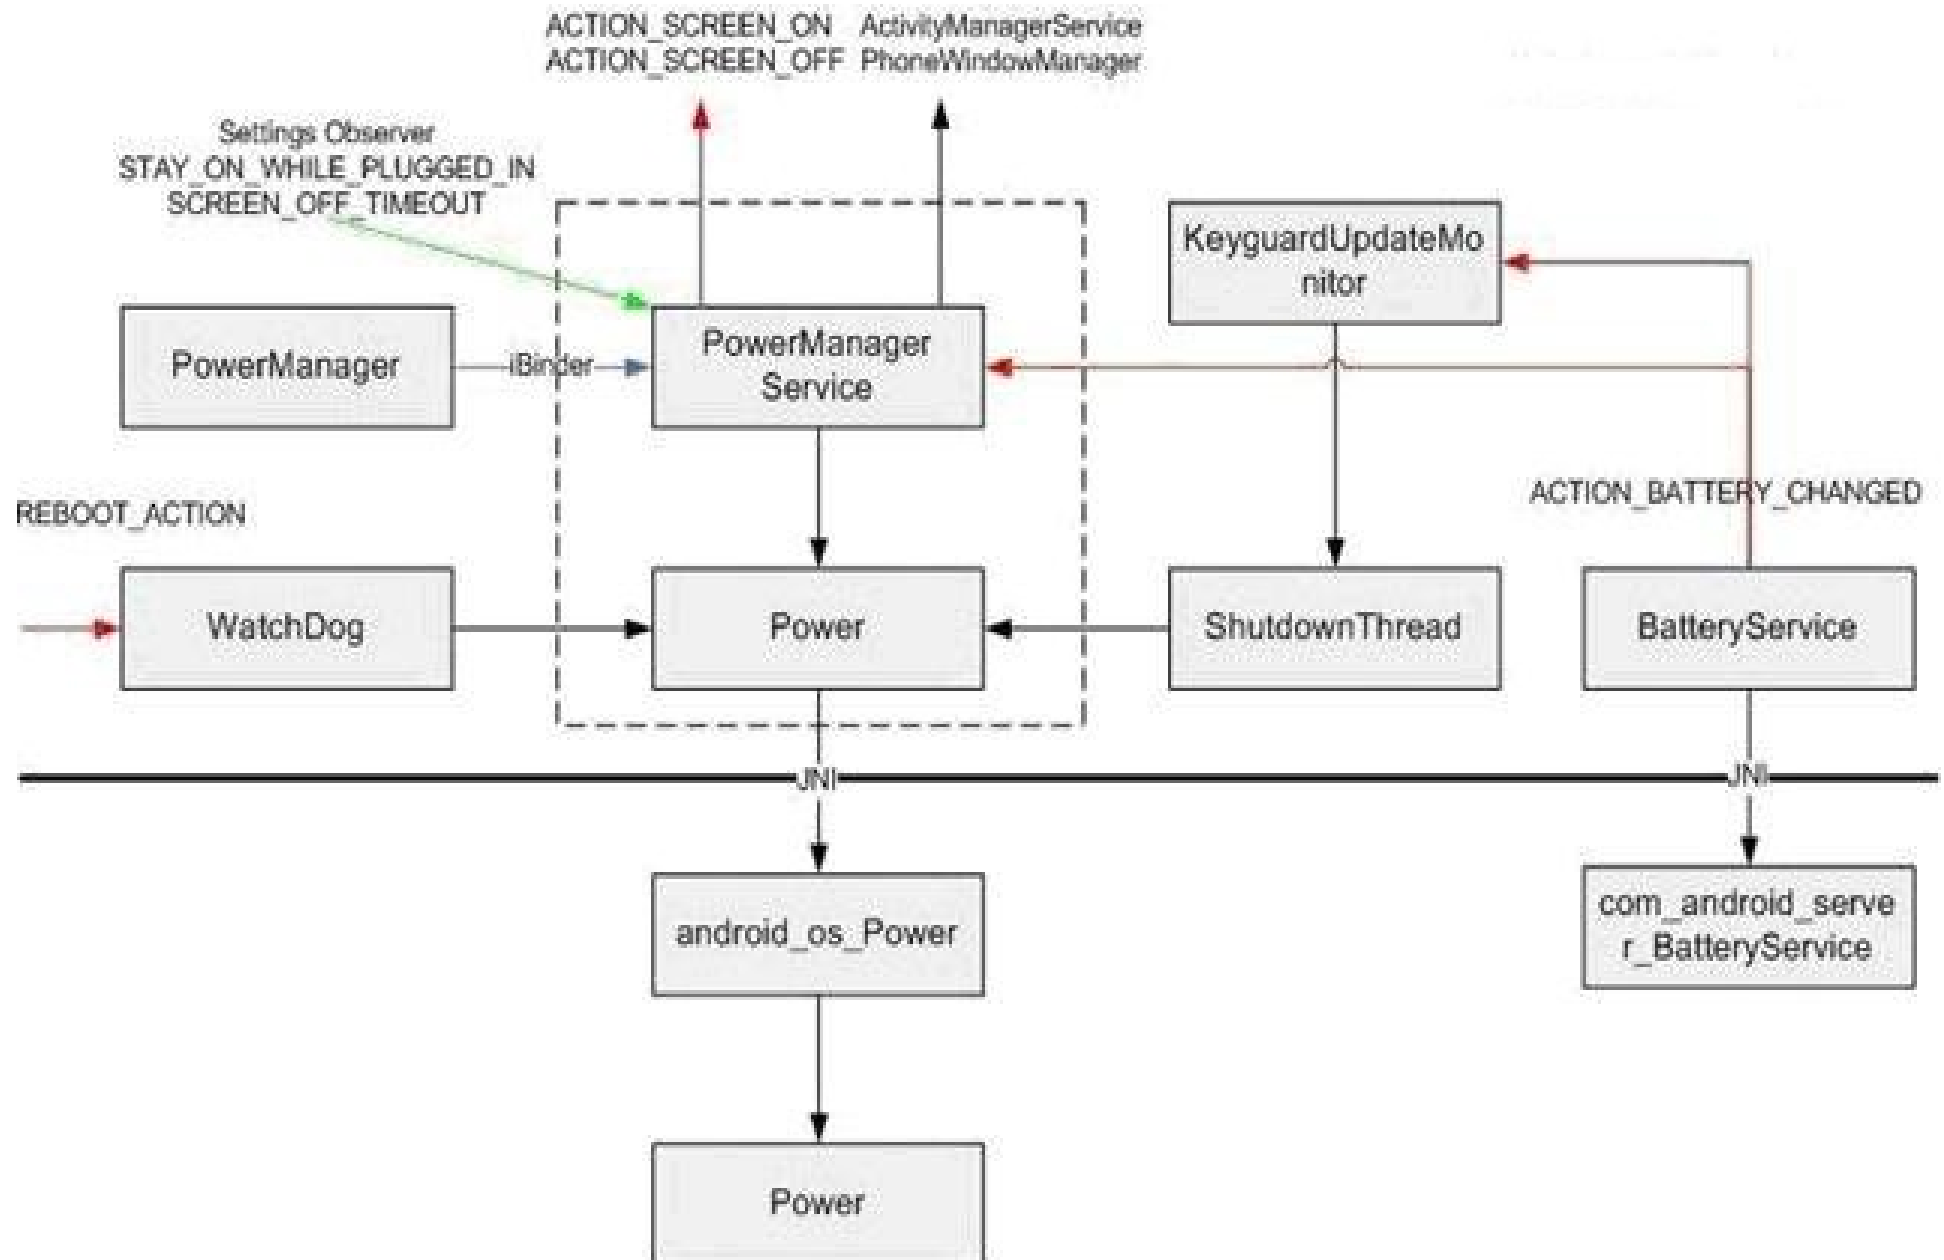
PM State Machine

Touchscreen or keyboard user activity event or full wake locks acquired.





- ▶ Introduction to Linux Power Management
- ▶ Concepts behind Android Power Management
 - ▶ Basically, it is the slim wrapper for Linux Power Management
- ▶ **Design and Implementation**
- ▶ Room for Improvements





android_os_Power

Settings Observe
STAY ON WHILE PLUG
SCREEN_OFF TIME

PowerManager

REBOOT_ACTION

WatchDog

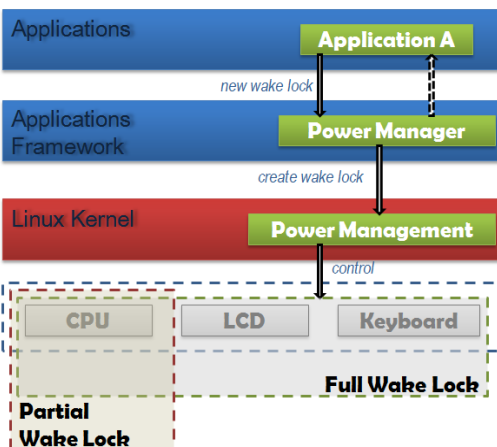
frameworks/base/core/jni/android_os_Power.cpp

```
...
static JNINativeMethod method_table[] = {
    { "acquireWakeLock", "(Ljava/lang/String;)V", (void*)acquireWakeLock },
    { "releaseWakeLock", "(Ljava/lang/String;)V", (void*)releaseWakeLock },
    { "setLastUserActivityTimeout", "(J)I", (void*)setLastUserActivityTimeout },
    { "setLightBrightness", "(II)I", (void*)setLightBrightness },
    { "setScreenState", "(Z)I", (void*)setScreenState },
    { "shutdown", "()V", (void*)android_os_Power_shutdown },
    { "reboot", "(Ljava/lang/String;)V", (void*)android_os_Power_reboot },
};

int register_android_os_Power(JNIEnv *env)
{
    return AndroidRuntime::registerNativeMethods(
        env, "android/os/Power",
        method_table, NELEM(method_table));
}
```

JNI

JNI



android_os_Power

Power

```
static void
acquireWakeLock(JNIEnv *env, jobject clazz,
                jint lock, jstring idObj)
{
    if (idObj == NULL) {
        throw_NullPointerException(env, "id is null");
        return ;
    }

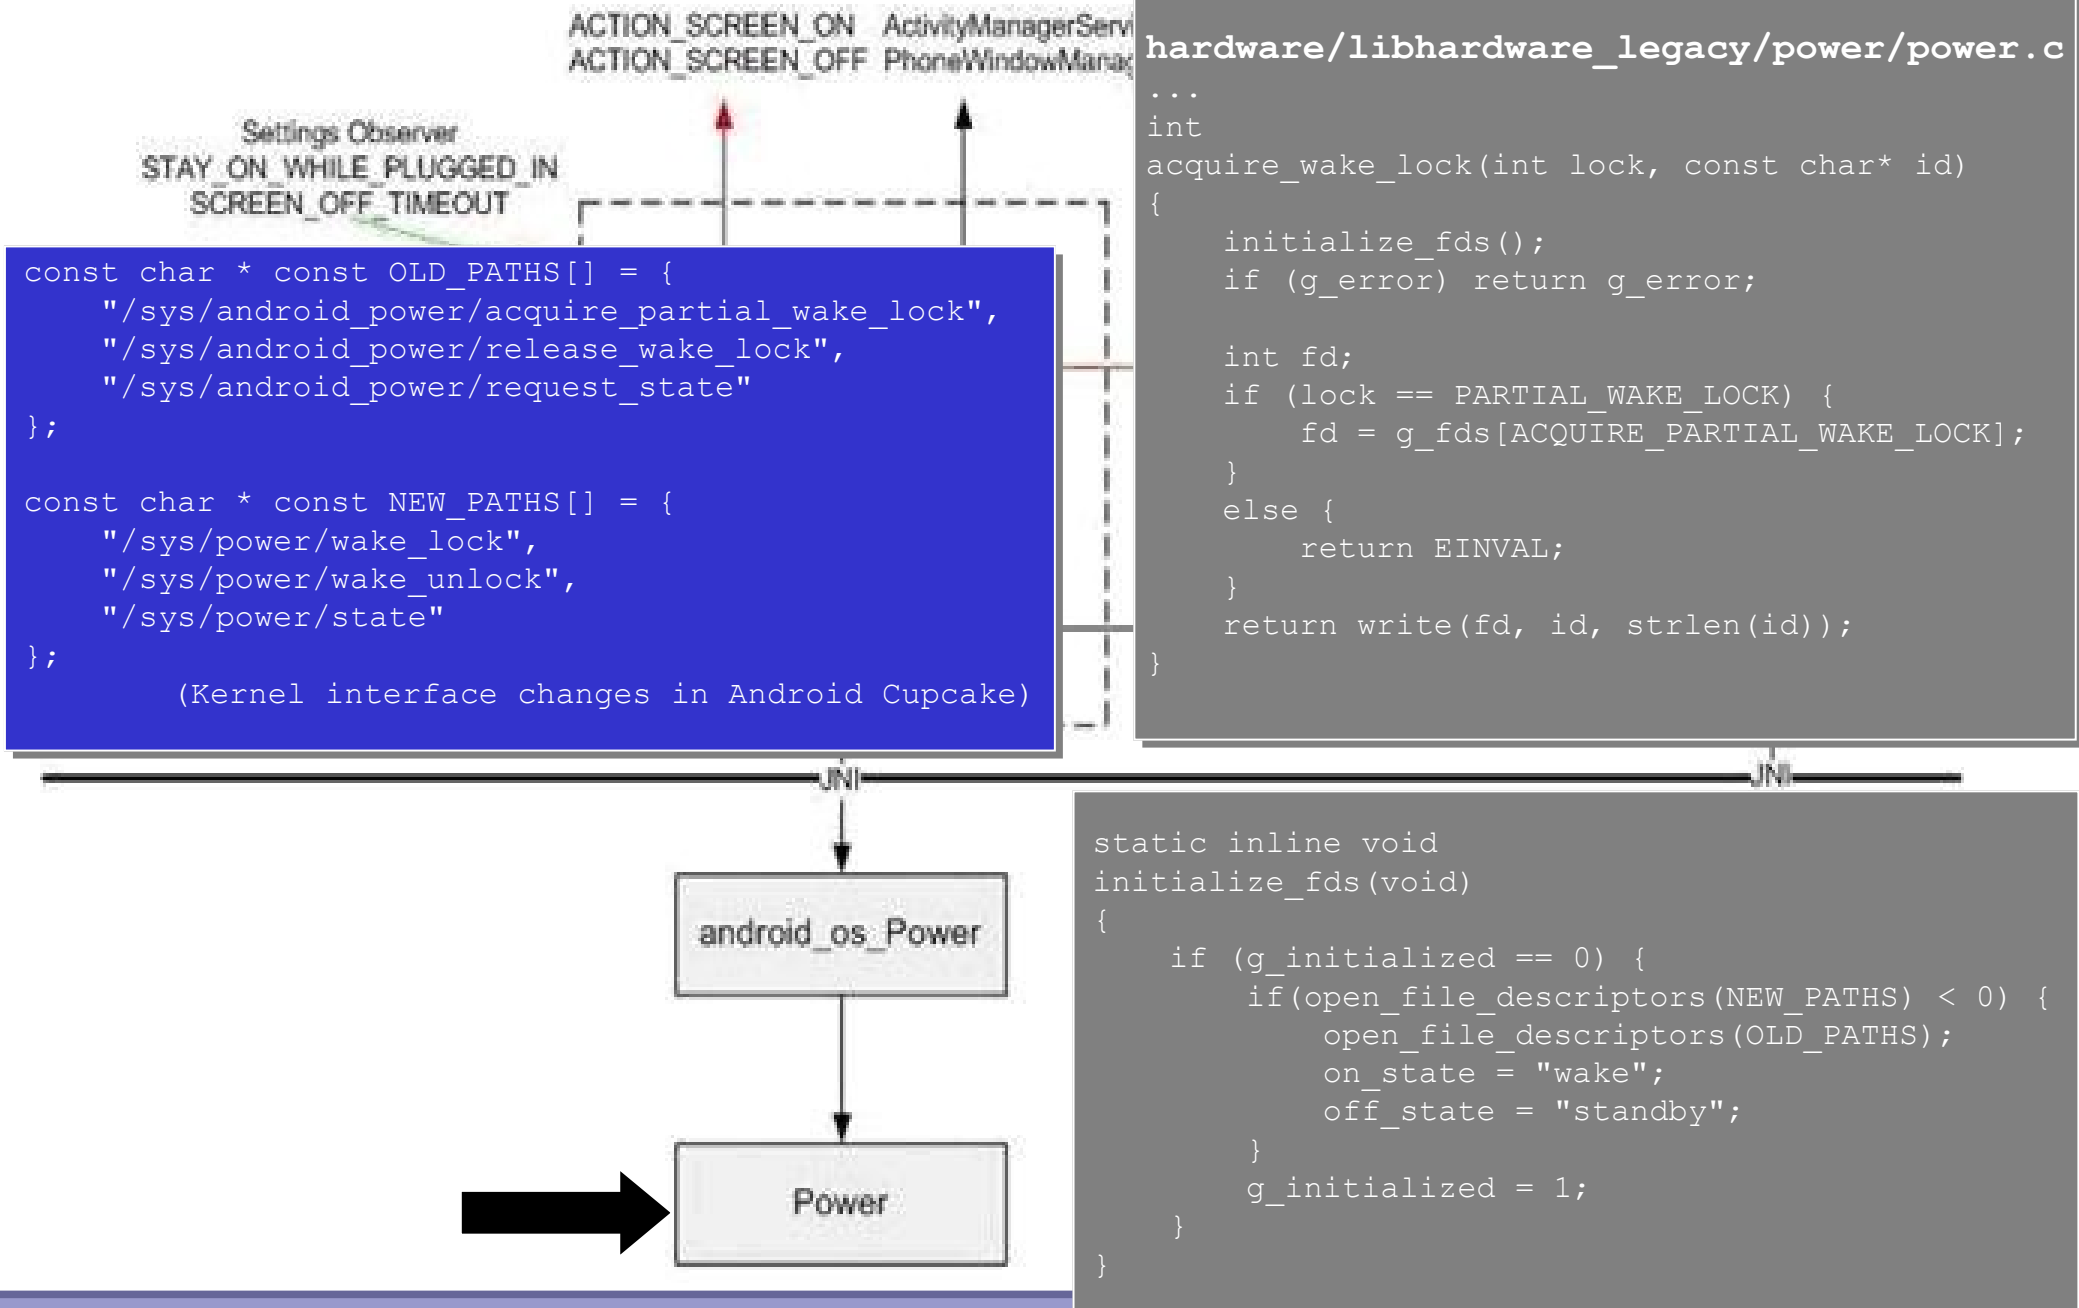
    const char *id = env->GetStringUTFChars(idObj, NULL);

    acquire_wake_lock(lock, id);

    env->ReleaseStringUTFChars(idObj, id);
}
```



Power





Android PM Kernel APIs

▶ Source code (Cupcake, linux-2.6.27)

▶ kernel/kernel/power/userwake.c

▶ kernel/kernel/power/wakelock.c

```
static int power_suspend_late(
    struct platform_device *pdev,
    pm_message_t state)
{
    int ret =
        has_wake_lock(WAKE_LOCK_SUSPEND) ?
        -EAGAIN : 0;
    return ret;
}

static struct platform_driver power_driver = {
    .driver.name = "power",
    .suspend_late = power_suspend_late,
};

static struct platform_device power_device = {
    .name = "power",
};
```

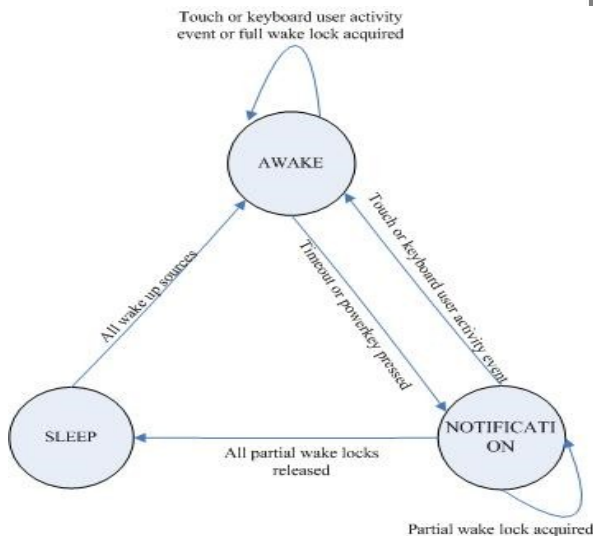
```
static long has_wake_lock_locked(int type)
{
    struct wake_lock *lock, *n;
    long max_timeout = 0;
    BUG_ON(type >= WAKE_LOCK_TYPE_COUNT);
    list_for_each_entry_safe(lock, n,
        &active_wake_locks[type], link) {
        if (lock->flags & WAKE_LOCK_AUTO_EXPIRE) {
            long timeout = lock->expires - jiffies;
            if (timeout <= 0)
                expire_wake_lock(lock);
            else if (timeout > max_timeout)
                max_timeout = timeout;
        } else
            return -1;
    }
    return max_timeout;
}

long has_wake_lock(int type)
{
    long ret;
    unsigned long irqflags;
    spin_lock_irqsave(&list_lock, irqflags);
    ret = has_wake_lock_locked(type);
    spin_unlock_irqrestore(&list_lock, irqflags);
    return ret;
}
```




Android PM Kernel APIs

▶ kernel/kernel/power/wakelock.c



```
static int __init wakelocks_init(void)
{
    int ret;
    int i;

    for (i = 0; i < ARRAY_SIZE(active_wake_locks); i++)
        INIT_LIST_HEAD(&active_wake_locks[i]);

    wake_lock_init(&main_wake_lock, WAKE_LOCK_SUSPEND, "main");
    wake_lock(&main_wake_lock);
    wake_lock_init(&unknown_wakeup, WAKE_LOCK_SUSPEND, "unknown_wakeups");


    ret = platform_device_register(&power_device);
    if (ret) {
        pr_err("wakelocks_init: platform_device_register failed\n");
        goto err_platform_device_register;
    }
    ret = platform_driver_register(&power_driver);
    if (ret) {
        pr_err("wakelocks_init: platform_driver_register failed\n");
        goto err_platform_driver_register;
    }

    suspend_work_queue = create_singlethread_workqueue("suspend");
    if (suspend_work_queue == NULL) {
        ret = -ENOMEM;
        goto err_suspend_work_queue;
    }
}
```



Back to userspace

- ▶ Use “`grep -r acquire_wake_lock`” to discover.
 - ▶ `frameworks/base/libs/ui/EventHub.cpp`
 - ▶ `EventHub::EventHub()`, `EventHub::~~EventHub()`,
`EventHub::getEvent()`
 - ▶ `hardware/ril/libril/ril.cpp`
 - ▶ `RIL_onUnsolicitedResponse()`,
 - ▶ `system/wlan/ti/sta_dk_4_0_4_32/CUDK/ti wlan_loader/ti wlan_loader.c`
 - ▶ `start_cli()`





Back to userspace

- ▶ Use “**grep -ri power**” to discover.
 - ▶ base/core/jni/android_net_wifi_Wifi.cpp
 - ▶ android.net.wifi.WifiNative.setPowerModeCommand
(android_net_wifi_setPowerModeCommand)
 - ▶ base/core/java/com/android/internal/os/BatteryStats.java
 - ▶ base/core/java/com/android/internal/os/BatteryStatsImpl.java
 - ▶ base/core/java/com/android/os/PowerManager.java
 - ▶ base/core/java/com/android/os/LocalPowerManager.java
 - ▶ base/core/java/com/android/os/Power.java
 - ▶ base/core/java/com/android/app/ApplicationContext.java
 - ▶ base/core/java/android/bluetooth/ScoSocket.java
 - ▶ base/core/java/android/bluetooth/HeadsetBase.java
 - ▶ base/core/media/java/android/media/MediaPlayer.java
 - ▶ base/core/media/java/android/media/AsyncPlayer.java



Back to userspace

- ▶ Use “**grep -ri power**” to discover.
 - ▶ base/telephony/java/com/android/internal/telephony/gsm/GSMConnection.java
 - ▶ base/telephony/java/com/android/internal/telephony/gsm/RIL.java
 - ▶ base/telephony/java/com/android/internal/telephony/gsm/GSMPhone.java
 - ▶ base/services/jni/com_android_server_BatteryService.cpp
 - ▶ base/services/java/com/android/server/am/ActivityManagerService.java
 - ▶ base/services/java/com/android/server/SystemServer.java
 - ▶ base/services/java/com/android/server/BatteryService.java
 - ▶ base/services/java/com/android/server/AlarmManagerService.java
 - ▶ base/services/java/com/android/server/LocationManagerService.java
 - ▶ base/services/java/com/android/server/HardwareService.java
 - ▶ base/services/java/com/android/server/PowerManagerService.java



Back to userspace

- ▶ Use “**grep -ri power**” to discover.
 - ▶ system/core/toolbox/powerd.c
 - ▶ system/core/toolbox/alarm.c
 - ▶ system/core/vold/mmc.c
 - ▶ system/core/vold/uevent.c

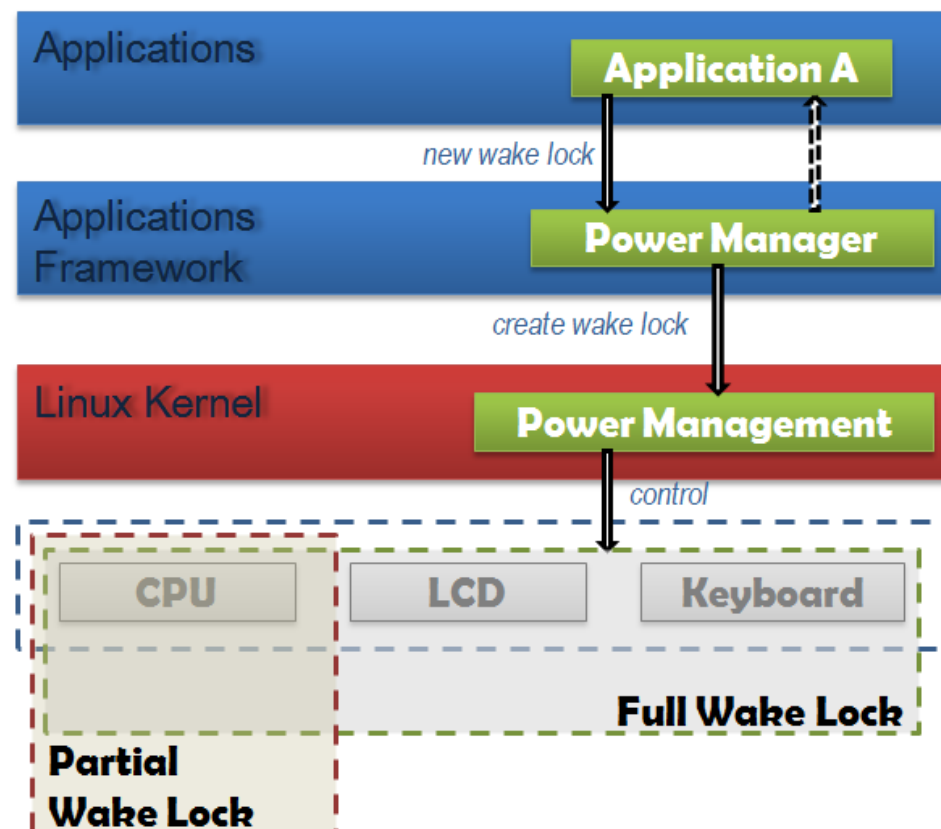


- ▶ Introduction to Linux Power Management
- ▶ Concepts behind Android Power Management
 - ▶ Basically, it is the slim wrapper for Linux Power Management
- ▶ Design and Implementation
- ▶ **Room for Improvements**

So, Android PM is simple

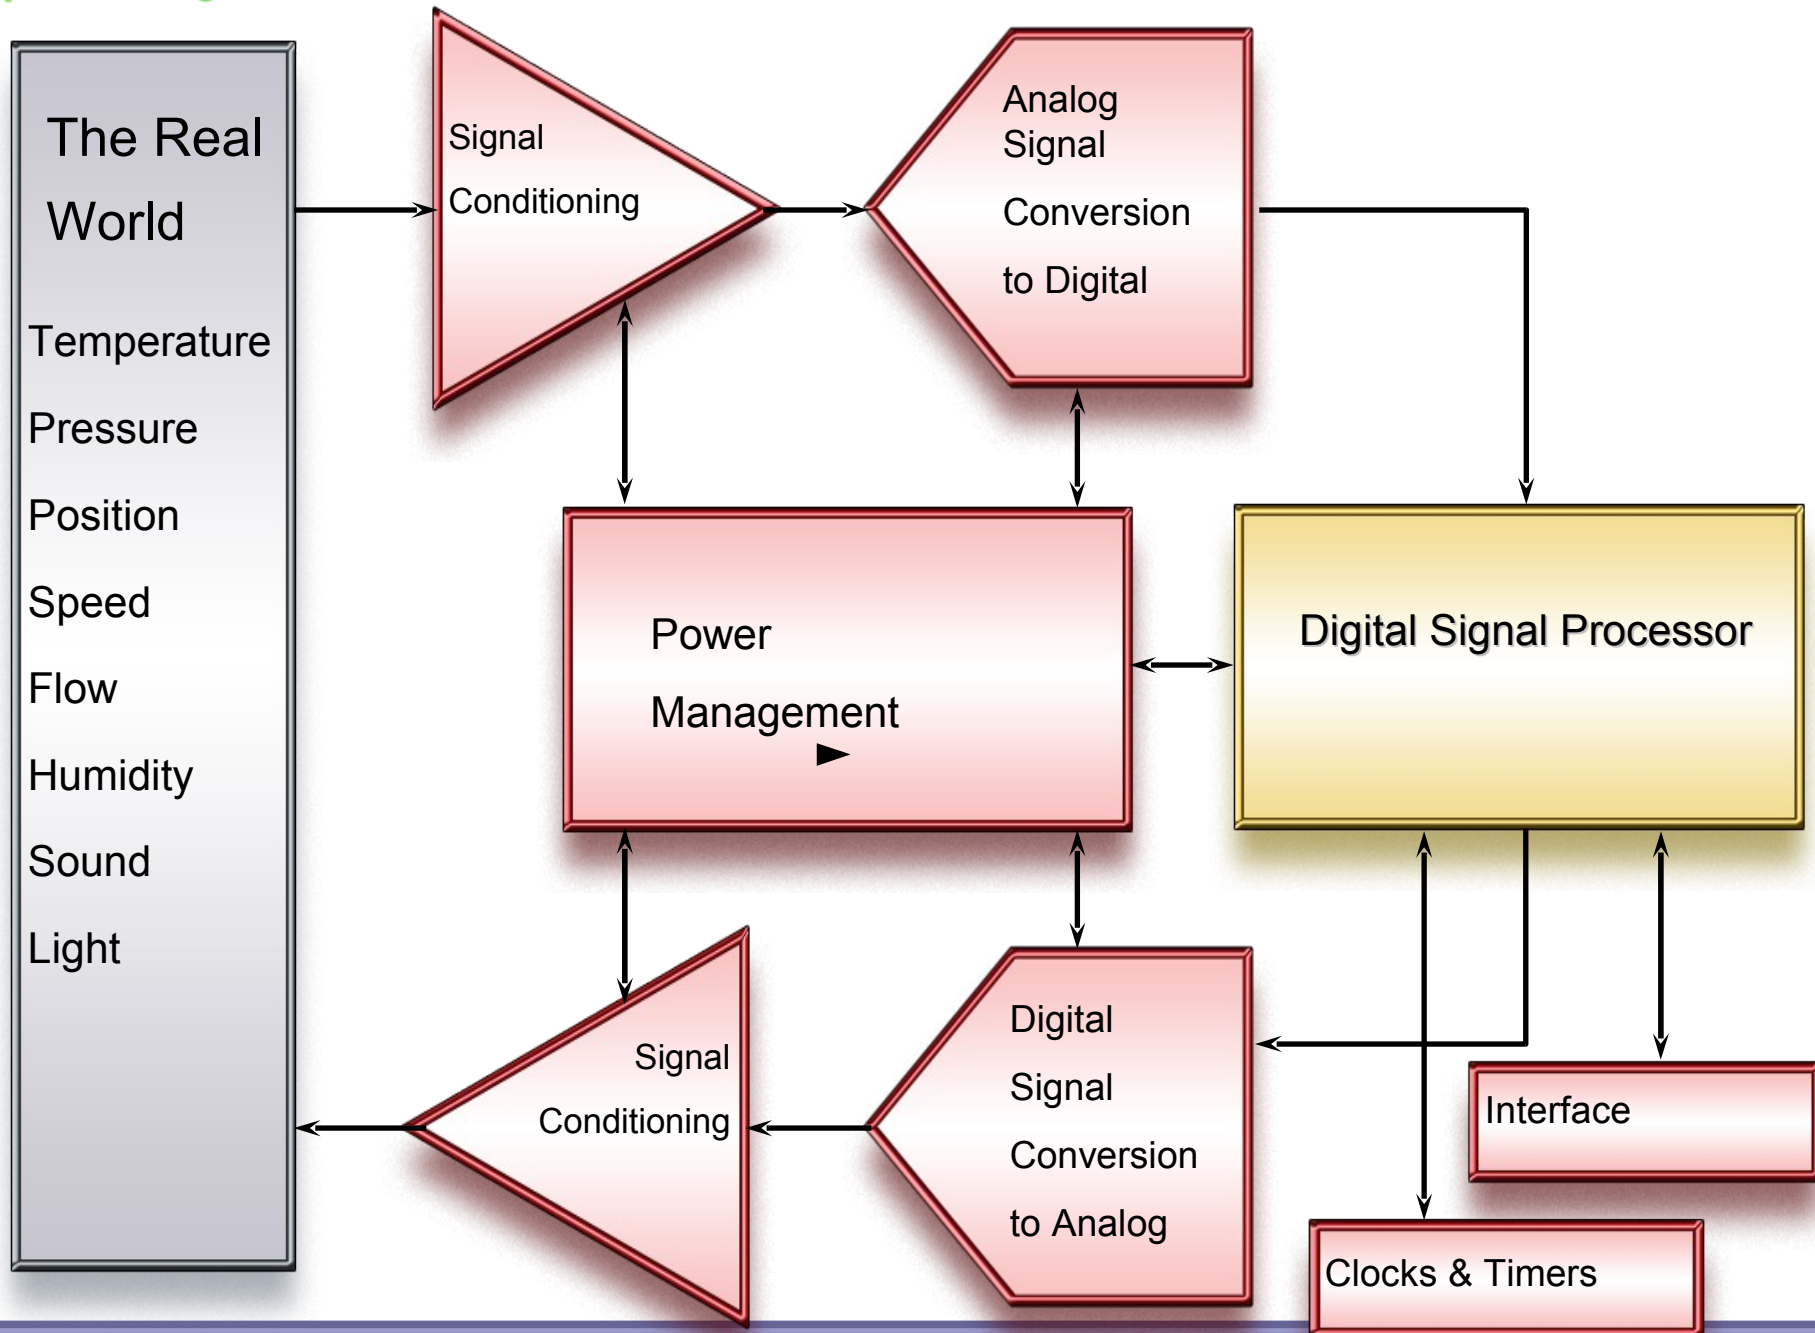
- ▶ Key concept is “Wake Locks”, which is simple and portable.
- ▶ The implementation should be introspected.

- ▶ If there are no active wake locks, CPU will be turned off.
- ▶ If there is are partial wake locks, screen and keyboard will be turned off.



- ▶ Nowadays, CPU can enter into more states for power saving and usability purpose! → applied in modern SoC

Interface to Physical World





Advanced Power Management

Power management aims to improve battery life of equipment by minimizing power consumption while guaranteeing expected system performance

Active power consumption occurs while some processing is on-going

- ▶ Dynamic power consumption (transistor switching) + Leakage consumption

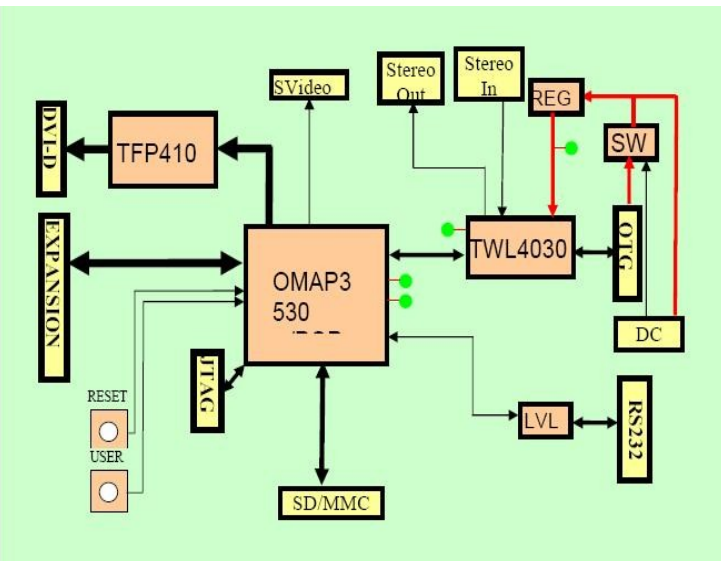
Static (also Standby or Idle) power consumption occurs when limited or no processing is on-going and the system is waiting for a wakeup event

- ▶ Very limited dynamic power consumption + Leakage consumption
- ▶ Managed by
 - ▶ Dynamic Voltage & Frequency Scaling (DVFS)
 - ▶ Adaptive Voltage Scaling (AVS)
 - ▶ Dynamic Power Switching (DPS)


On OMAP35xx, power management is handled by the Power, Reset and Clock Management (PRCM) module

Adaptive Voltage Scaling

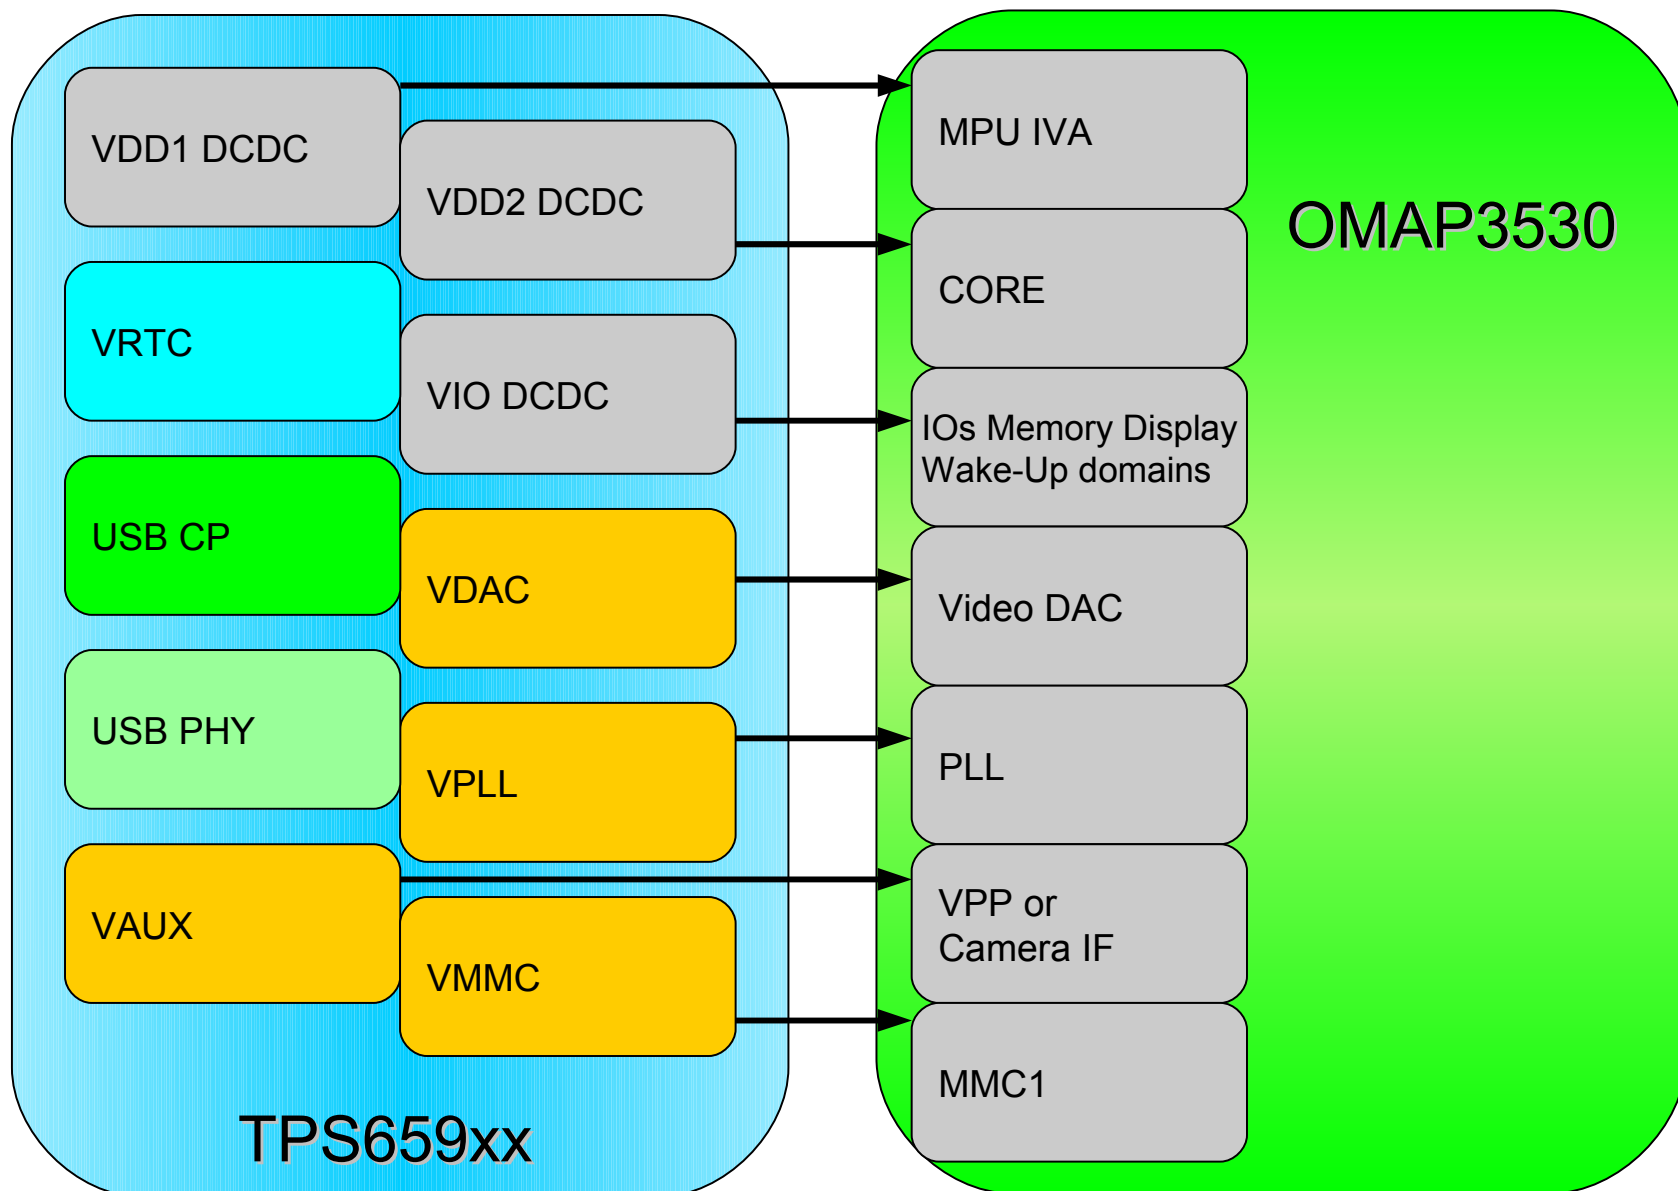
- ▶ Silicon manufacturing process yields a distribution of performance capability
- ▶ For a given frequency requirement:
 - ▶ Devices on hot/strong/fast end of distribution can meet this at a lower voltage
 - ▶ Devices on cold/weak/slow end of distribution need higher voltage
- ▶ Simple system will set the higher voltage for operating all devices
- ▶ Smarter system will adapt operating voltage per device.

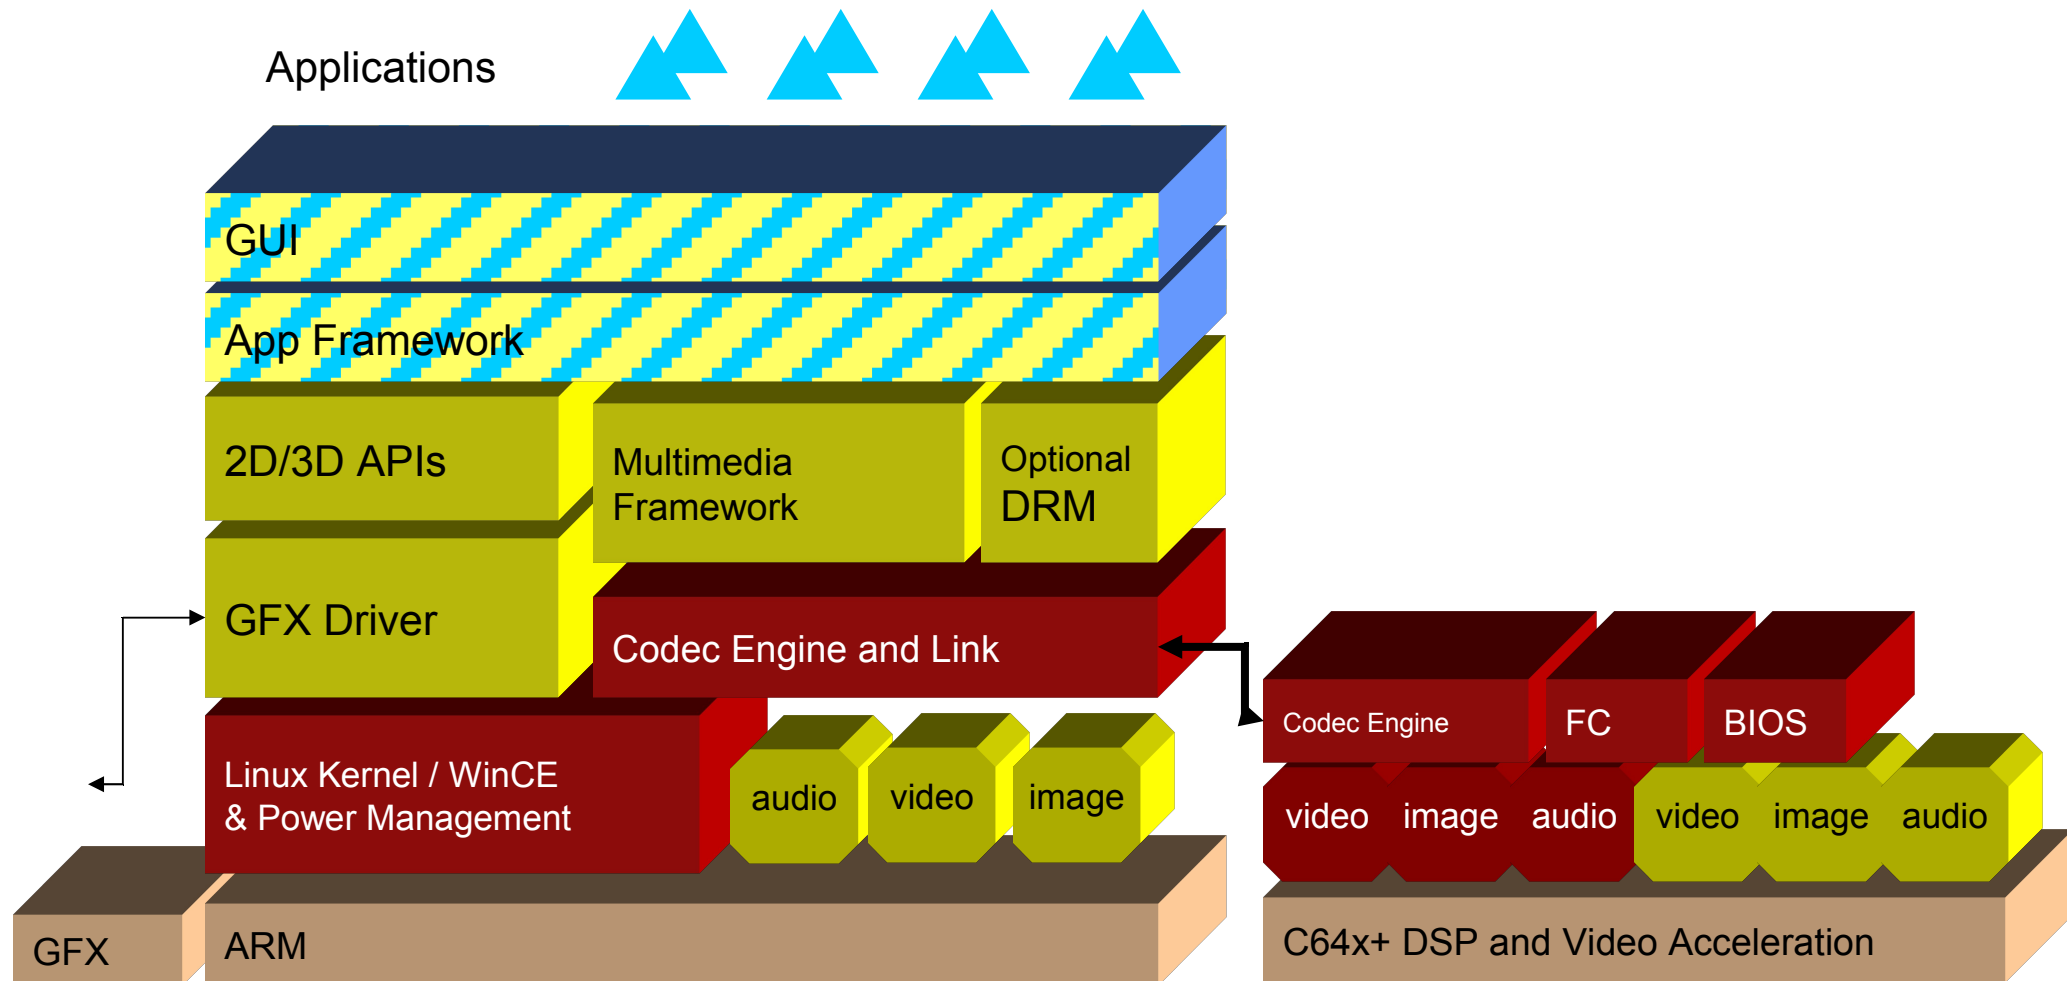


hotter device



Power Block Diagram







Linux PM Mechanisms

- ▶ cpuidle
 - ▶ Generic framework for supporting software-controlled idle processor power management
 - ▶ Hardware specific drivers
 - ▶ Various governing for the state transition decisions
- ▶ Latency and power management
 - ▶ Framework for expressing latency constraints, and make sure that they are taken into account for power management decisions
- ▶ Needs to be integrated into Android Partial Wake Locks

Integration is the key to the most power-saving system





- ▶ PDK :: Android – Power Management

- ▶ http://www.netmite.com/android/mydroid/development/pdk/docs/power_management.html

- ▶ Android Cupcake source code

- ▶ <http://www.netmite.com/android/mydroid/cupcake/>

- ▶ BeagleBoard and its Linux support

- ▶ <http://elinux.org/BeagleBoard>

- ▶ class PowerManager

- ▶ <http://developer.android.com/reference/android/os/PowerManager.html>

- ▶ Free-Electrons

- ▶ <http://free-electrons.com/training>

- ▶ CELF's power management specification

- ▶ http://elinux.org/Power_Management