

Power, Energy and Performance

Kunal Korgaonkar

Prof. Tajana Simunic Rosing

Department of Computer Science and Engineering

UCSD

Motivation for Power Management

- Power consumption is a critical issue in system design today
 - Mobile systems: maximize battery life
 - High performance systems: minimize operational costs




➔ 1.2% of total electrical use in US
devoted for powering and cooling
data centers

➔ \$1 operation => \$1 cooling

Power and Energy Relationship

$$E = \int P dt$$




Low Power vs. Low Energy

- Minimizing the **power consumption** is important for
 - the design of the power supply
 - the design of voltage regulators
 - the dimensioning of interconnect
 - short term cooling
- Minimizing the **energy consumption** is important due to
 - restricted availability of energy (mobile systems)
 - limited battery capacities (only slowly improving)
 - very high costs of energy (solar panels, in space)
 - cooling
 - high costs
 - limited space
 - dependability
 - long lifetimes, low temperatures

Intuition


- System and components are:
 - Designed to deliver *peak performance*, but ...
 - Not needing peak performance most of the time
- Dynamic Power Management (DPM)
 - Shut down components during idle times
- Dynamic Voltage Frequency Scaling (DVFS)
 - Reduce voltage and frequency of components
- System Level Power Management Policies
 - Manage devices with different power management capabilities
 - Understand tradeoff between DPM and DVFS

DPM/DVFS



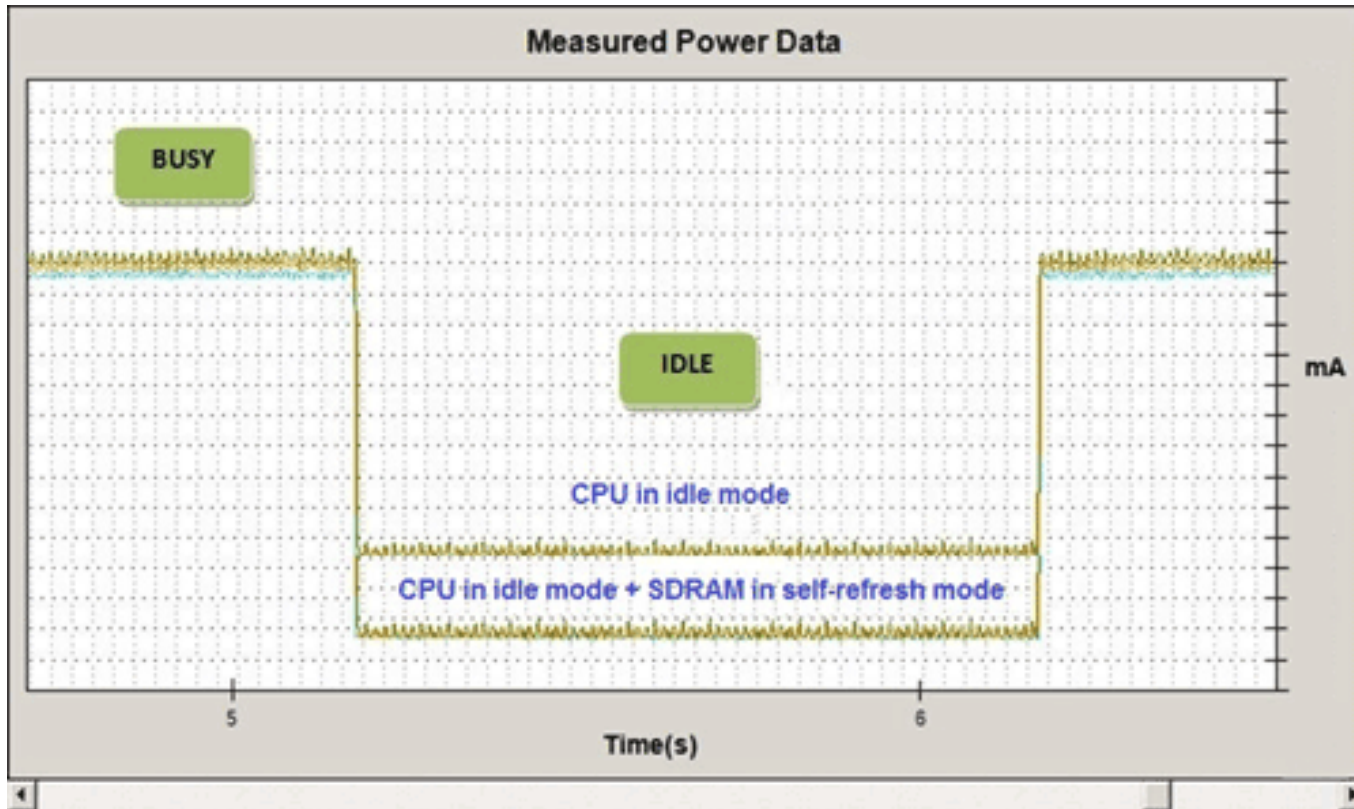
DVFS during active phases and
DPM during idle phases.

DVFS



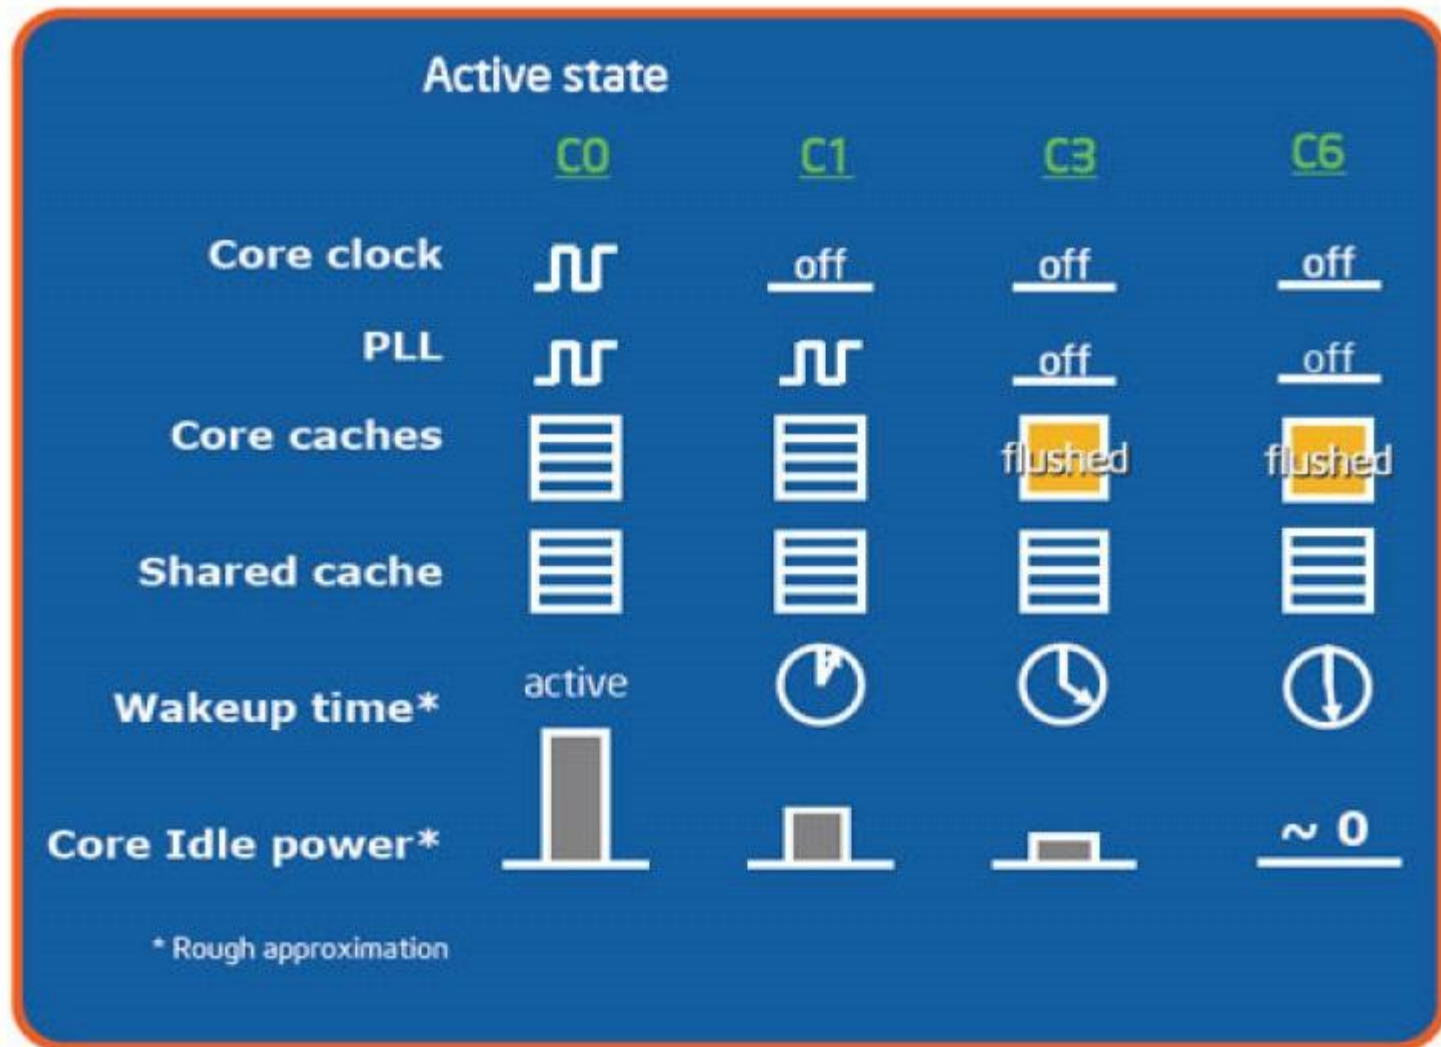
DVFS is an effective way of reducing the CPU energy consumption by providing “just-enough” computation power.

DPM



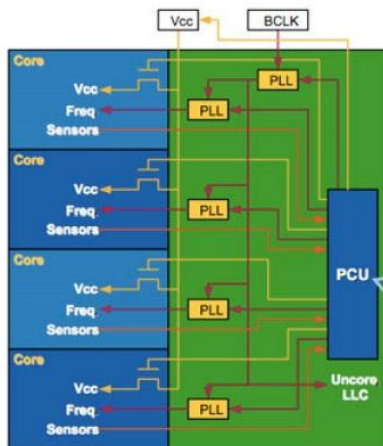
DPM is an effective way of reducing the CPU energy consumption by shutting down (completely or partially) when components are not in use.

Idle States



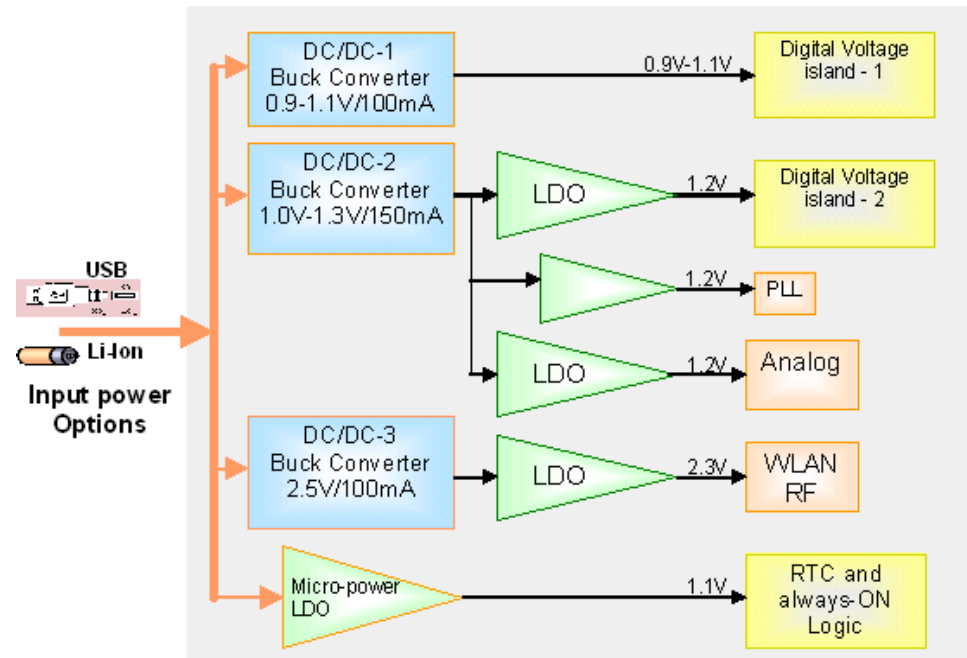
Power Supply

Power Control Unit



Integrated proprietary microcontroller
Shifts control from hardware to embedded firmware
Real time sensors for temperature, current, power
Flexibility enables sophisticated algorithms, tuned for current operating conditions


Figure 4: Nehalem's Power Control Unit (PCU)





Android/Linux Power Management and Class Projects

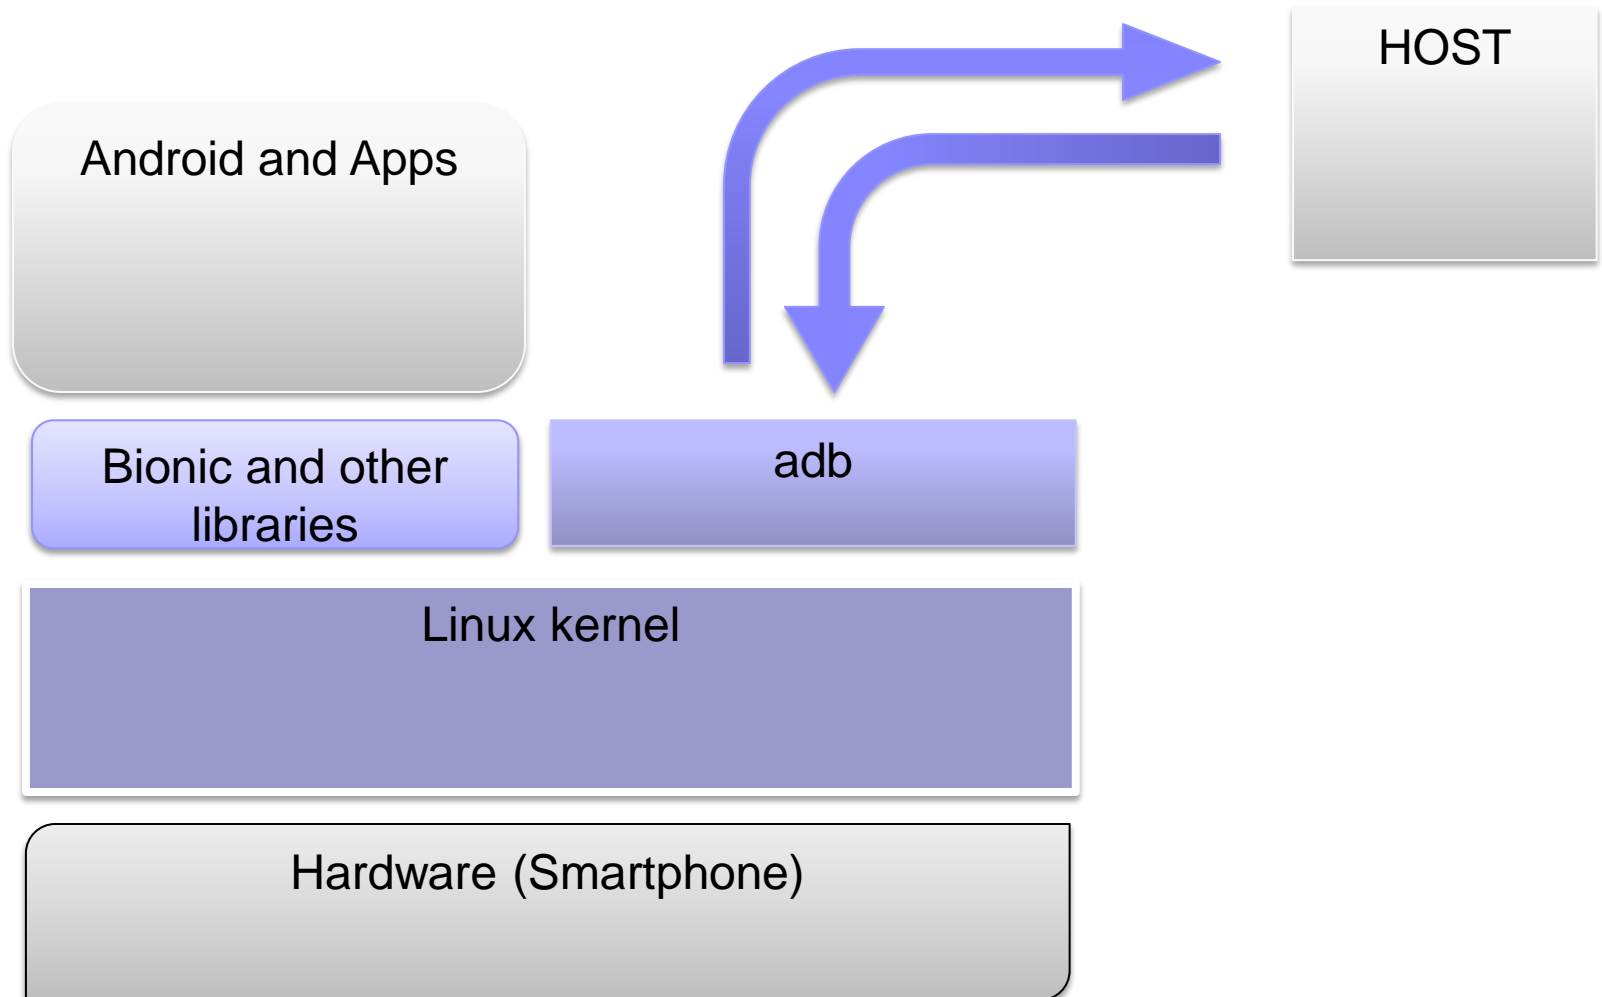
Frequency Governors



Project Part 2 – to help you get started

- Given a trace of cpu (IPC) and memory utilization (MPC) finding the best cpu frequency policy
- All on paper (no devices yet)
- Read about cpu frequency governors
 - http://docs.fedoraproject.org/en-US/Fedora/15/html/Power_Management_Guide/index.html
 - <http://doc.opensuse.org/documentation/html/openSUSE/opensuse-tuning/cha.tuning.power.html>

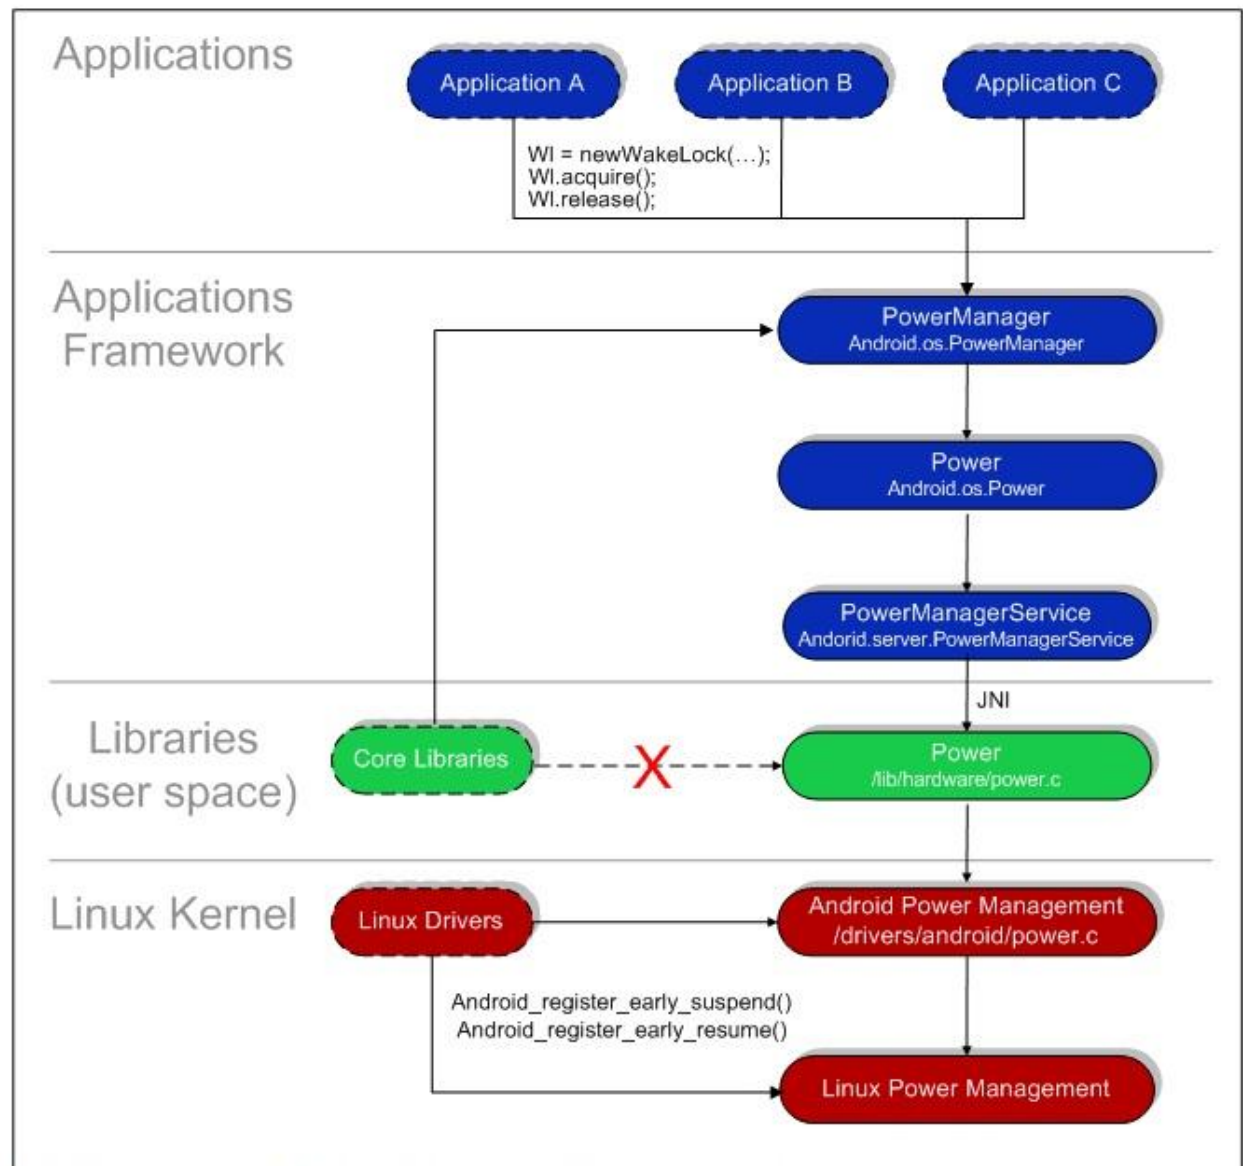
Experimental Setup



Android power management

- Android requires that applications and services request CPU (and other hardware resources) with "wake locks" through the Android application framework and native Linux libraries
- Examples
 - WAKE_LOCK_SUSPEND: prevents a full system suspend
 - WAKE_LOCK_IDLE: low-power states, which often cause large interrupt latencies or that disable a set of interrupts, will not be entered from idle until the wake locks are released
 - SCREEN_BRIGHT_WAKE_LOCK: Wake lock that ensures that the screen is on at full brightness; the keyboard backlight will be allowed to go off

Android Power Management



How it works

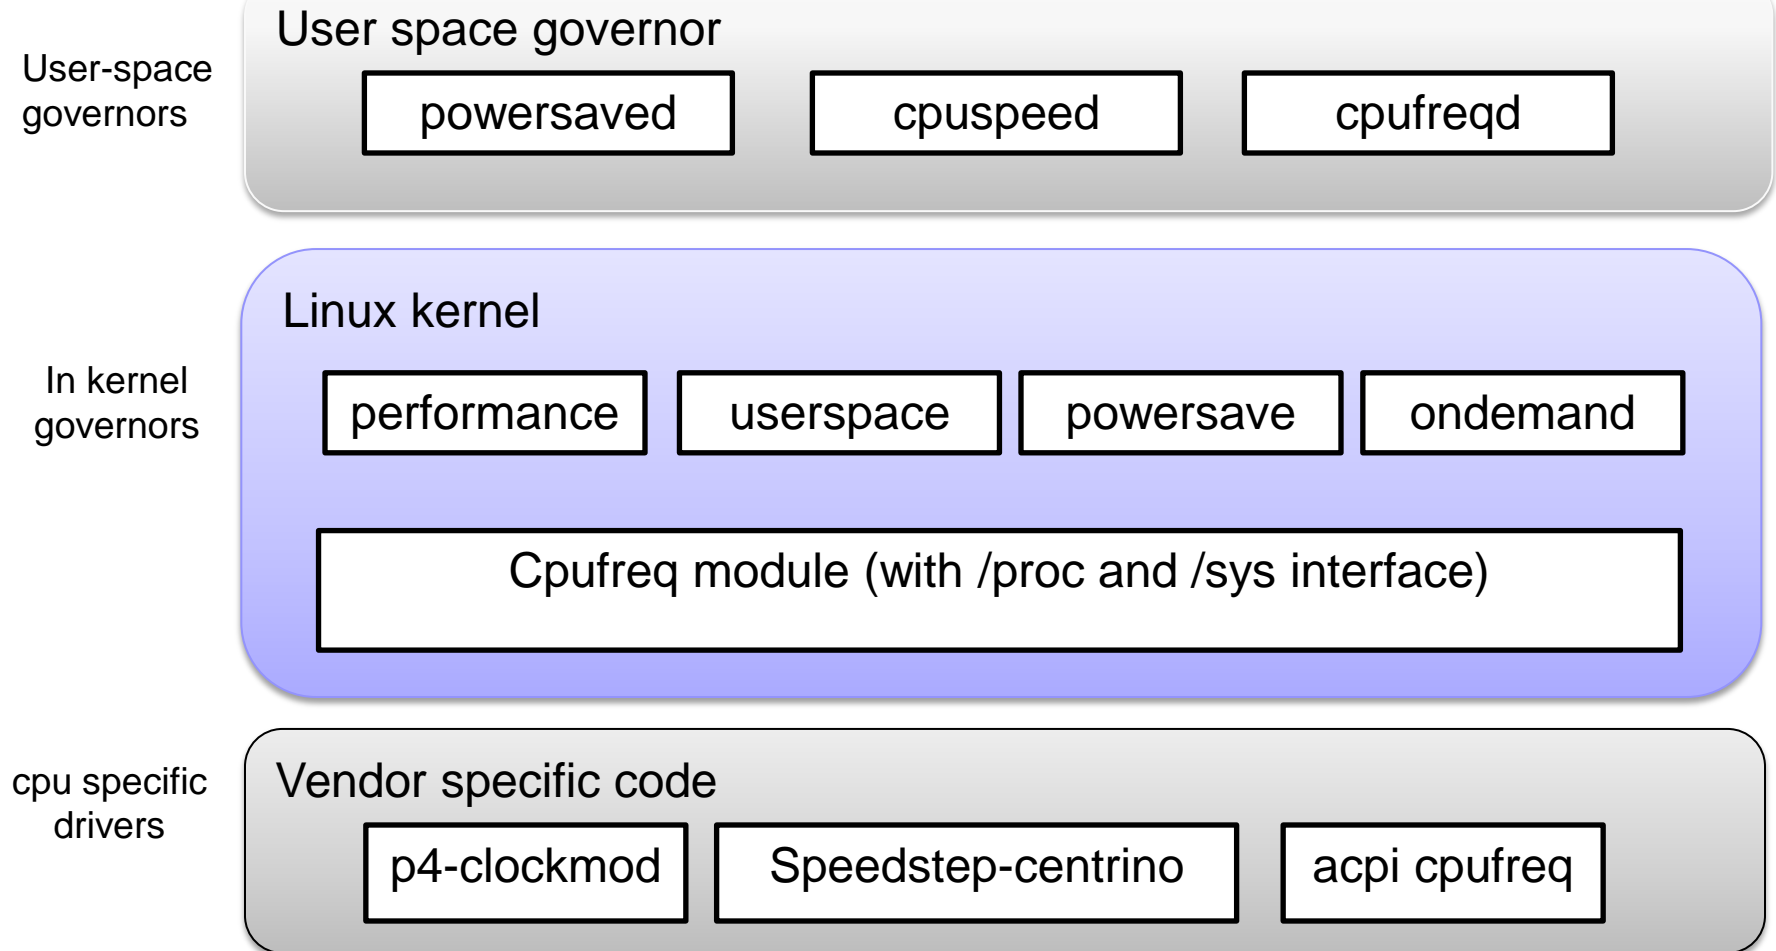
- **PowerManager class:**

- The Android Framework exposes power management to services and applications through the PowerManager class

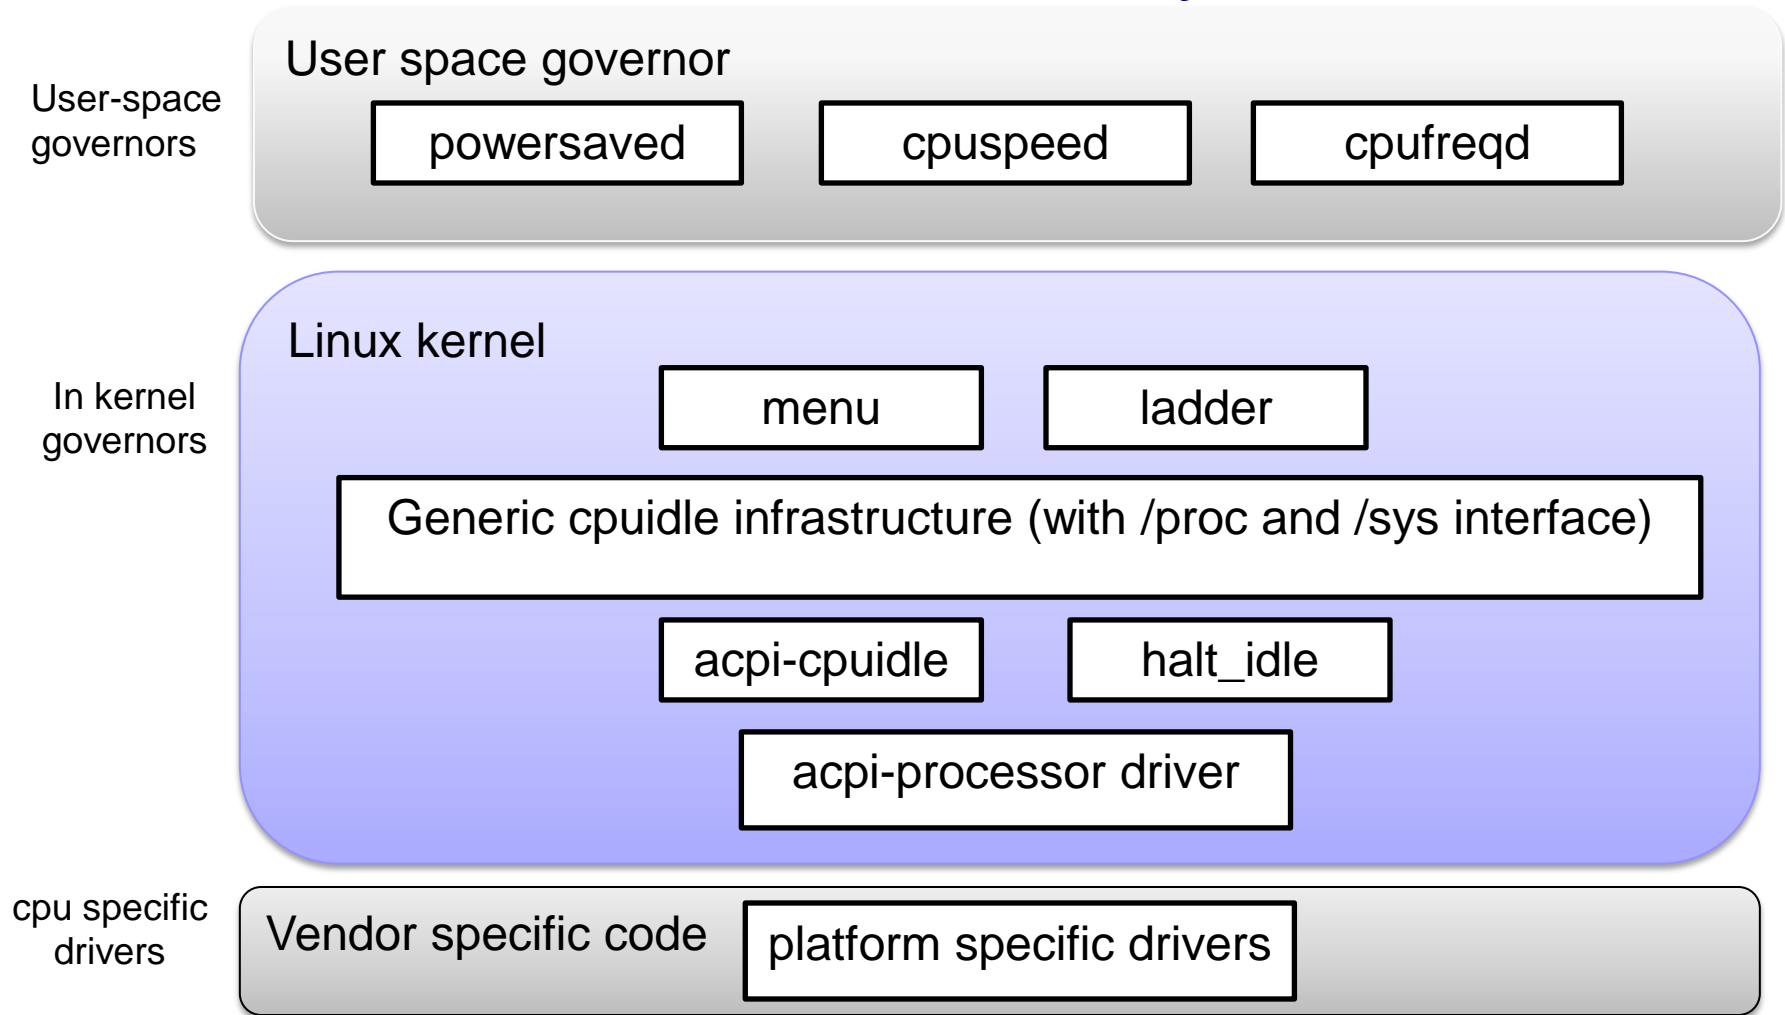
- **Registering Drivers with the PM Driver**

- You can register Kernel-level drivers with the Android Power Manager driver so that they're notified immediately before power down or after power up
 - E.g. Display driver can be registered to completely power down when a request comes in to power down from the user space

Linux CPU Frequency Subsystem



Linux Idle Subsystem



Workloads

- SPEC benchmarks

- ☐ Standardized applications
- ☐ But not mobile specific

- Apps

- ☐ Written with real market needs in mind
- ☐ But not standardized

- Micro-benchmarks

- ☐ Can be used to understand/study a specific mobile SOC subsystem
- ☐ Easy to debug

Future Project Activities (all in parallel)

Data collection

Data Analysis

Understanding and coming up with power
management policies

Project Activity: Data collection

- Session based
 - 2 to 5 mins
 - Multiple runs needed
 - For all three workloads
- Full 'working day' based
 - Full day logs
 - You will be the user

Part Activity: Data Analysis

- Use logs of real apps whenever possible
- As instructed use
 - SPEC benchmarks and
 - Micro-benchmarks



Project Activity: Understanding and coming up with your own power management policies

- Coming up with your own frequency governor
- TBA

What you will learn through these projects

- You will get to learn about
 - Android and Linux kernel
 - Mobile SOC hardware capability
 - Linux CPU frequency subsystem
 - How to use cpu frequency subsystem for power management
 - How benchmarks and real apps behave
 - Understanding performance, power and 'response time' trade-offs